

Why SAT-based Analysis of Large Real-world Feature Models is Easy

Jia Hui Liang
University of Waterloo
Ontario, Canada
jliang@gsd.uwaterloo.ca

Venkatesh Raman
Institute of Mathematical
Sciences, India
vraman@imsc.res.in

Vijay Ganesh
University of Waterloo
Ontario, Canada
vijay.ganesh@uwaterloo.ca

Krzysztof Czarnecki
University of Waterloo
Ontario, Canada
kczarnec@gsd.uwaterloo.ca

ABSTRACT

Modern Boolean SAT solvers are often employed in performing automatic analysis of real-world feature models (FM). These solver-based analyses are surprisingly very effective and scale well even though these real-world models are very large and these analysis problems are known to be NP-complete. To better understand why SAT solvers are so effective, we systematically studied many characteristics of real-world FMs and tested several distinct hypotheses. We eventually discovered that a key reason why large real-world FMs are easy-to-analyze is that a vast majority of the variables in these models are *unrestricted*, i.e., the models are satisfiable for both true and false assignments to such variables under suitable satisfiable partial assignments to its remaining variables. Once we made this discovery, it immediately followed that solvers can easily find satisfying assignments for such models without too many backtracks relative to the model-size, explaining why solvers scale so well. Further analysis led us to the conclusion that the presence of unrestricted variables in these real-world models can be attributed to their high-degree of *variability*. Leveraging these insights, we implemented a series of non-backtracking simplification techniques that are effective in eliminating most variables from such models. The variables/clauses that remain, called the “core”, tend to be very small that are easily solved. We also compared our findings with previous research that characterizes the difficulty of “randomly generated” FMs in terms of treewidth. Surprisingly, we found that treewidth is, at best, weakly correlated with the running time of modern SAT solvers for real-world FMs.

1. INTRODUCTION

Feature models (FM) are widely used to represent the variability and commonality in reusable software, first intro-

duced in 1990 [13]. Feature models define all the valid feature configurations of a product line, whether it be a car or a piece of software. The process of feature modeling helps to ensure that relevant features are reusable, and unused features are removed to lower the complexity of the system under design [9]. Promoting the reuse of features can shorten the time of delivering new products and lower the cost of production. Such models can automatically be analyzed using SAT solvers and model-checkers to detect inconsistencies or errors in the design process of a product, thus lowering the cost of production.

Modern conflict-driven clause learning (CDCL) Boolean SAT solvers (or simply SAT solvers¹) are known to efficiently solve many large real-world SAT instances obtained from a variety of domains such as software testing, program analysis and hardware verification. More recently, inspired by the success of SAT solvers in other domains, many researchers proposed the use of solvers to analyze feature models [2, 3, 4]. Subsequently, solvers have been widely applied to perform all manner of analysis on feature models, and have proven to be surprisingly effective even though the kind of feature model analysis discussed here is known to be an NP-hard problem², and real-world FMs tend to be very large, often running into hundreds of thousands of clauses and tens of thousands of variables. This state of affairs has perplexed practitioners and theoreticians alike.

Problem Statement: Hence, the problem we address in this paper is “Why are SAT-based analyses, that use modern CDCL SAT solvers, so effective at analyzing large real-world FMs?” The SAT-based analysis we are concerned about in this paper is the following: Given an FM, does it encode a valid product configuration. In SAT parlance, this corresponds to the question “Given a Boolean formula (equivalently FM) does it have a satisfying assignment (valid product configuration)?”

Our Findings: What we found was that the efficiency of

¹While there are other kinds of Boolean SAT solvers, such as ones based on Belief propagation, conflict-driven clause-learning are the dominant variety.

²This is easily shown by way of a polynomial-time reduction that encodes arbitrary Boolean 3-CNF formulas as feature models.

modern SAT solvers, at least in the case of real-world FMs, can be explained through a systematic and scientific analysis of the structure of these FMs. In particular, we found that the overwhelming number of variables occurring in real-world FMs (or its equivalent Boolean formula) that we studied are *unrestricted*. We say that a variable v in a Boolean formula ϕ given partial assignment S is unrestricted if there exist two satisfying extension of S (i.e., extensions of the partial assignment of S to all variables in ϕ such that ϕ is satisfiable), one with $v = \text{true}$ and the other with $v = \text{false}$. Intuitively, unrestricted variables do not cause modern CDCL SAT solvers to backtrack, simply because, under suitable partial assignments, the solver cannot make a mistake in assigning such variables the value *true* or *false*. This is important because if the number of backtracks a solver performs remains small relative to the increasing size of inputs, then the solver is likely to scale very well for such inputs. Indeed, this is exactly what we observed in our experiments, i.e., CDCL SAT solvers perform very few (near-constant) number of backtracks even as the size of FMs increase. We noted that this is directly attributable to unrestricted variables in such FMs.

The next question we addressed was “where do these unrestricted variables in FMs come from?”. We observed that the large percentages of unrestricted variables in real-world FMs is attributable to the *high variability* in such models. We say that an FM has high variability, if a large number of features occurring in such a model are optional, i.e., one can obtain a valid product configuration irrespective of whether such features are selected. Indeed, this observation is consistent with the fact that FMs tend to capture all features of a product line, whereas deriving a valid product from such models only require relatively few features to be present (resp. absent).

We also noted that real-world FMs have shallow but broad trees and many leaves are *optional features*. These features tend to map to unrestricted variables because they are rarely forced to be present or absent. Optional leaf features, if unconstrained by cross-tree constraints, will appear once in the propositional formula and hence corresponds to *pure variables*³. Similarly, if the feature is very lightly constrained, the corresponding Boolean variable might be *almost pure*, e.g., most occurrences of the variable in the formula have the same sign. It turns out such variables are very amenable to many well-known Boolean logic simplifications such subsumption, pure literal rule and general variable elimination. We leveraged this observation to implement a set of standard simplifications that for the most part completely solved large real-world FMs without ever needing to call a SAT solver.

Contributions: Here we briefly describe the contributions made in this paper.

1. As described above, we provide an explanation for why SAT solvers are effective and scalable for large real-world FMs. The explanation is that the overwhelming majority of the variables in real-world FMs are *unrestricted*, and solvers tend not to backtrack in the pres-

³A variable x in a formula ϕ is called pure, if it only occurs positively or negatively

ence of such variables. We argue that the reason for the presence of large number of unrestricted variables in real-world FMs has to do with high variability in such models. We also found that if we switch off all the heuristics in modern SAT solvers, except Boolean constant propagation (BCP) and backjumping (no clause-learning), then the solver does not suffer any deterioration in performance while solving FMs. This observation is consistent with the fact that the vast majority of variables in real-world FMs are unrestricted that are easily solved through unit resolution (BCP) and minimal backjumping.

2. Based on the above findings we implemented a number of well-known non-backtracking, albeit incomplete, simplification routines that are very effective in solving large real-world FMs. Often these simplifications completely solve such instances. There were few instances where these simplifications didn’t completely solve the input FMs, and instead returned a very small simplified formula, we call a “core”. These cores tend to consist largely of *Horn* clauses, that are subsequently easily solved by modern CDCL solvers.

The point of these experiments was to reinforce that most of the variables occurring in real-world FMs are *unrestricted*, which are easily eliminated by the class of simplifications discussed in the paper.

3. Following previous work by Pohl et al. [22], we performed experiments to see if the treewidth of graphs of the formulas obtained from FMs correlates with time. We found that for large real-world FMs the correlation is weak.
4. We also developed a technique for generating hard artificial FMs, to better understand the differences between the characteristics of large real-world, but easy, FMs vs. hard artificial ones.

2. BACKGROUND

This section provides the necessary background on FMs and the use of SAT solvers in analyzing them.

2.1 Feature Models

Structurally, a feature model looks like a tree where each node represents a distinct feature. (The terms node and feature will be used interchangeably.) Child nodes have two flavours: *mandatory* (the child feature is present if and only if the parent feature is present) and *optional* (if the parent feature is present then the child feature is optional, otherwise it is absent). Parent nodes can also restrict its children with feature groups: *or* (if the parent feature is present, then at least one of its child features is present) and *alternative* (if the parent feature is present, then exactly one of its child features is present). These are the structural constraints between the child and the parent. Figure 1 is an example of a feature model represented diagrammatically.

Structural constraints are often not enough to enforce the integrity of the model, in which case cross-tree constraints are necessary. Cross-tree constraints have no restrictions like structural constraints do, and can apply to any feature regardless of their position in the tree-part of the model.

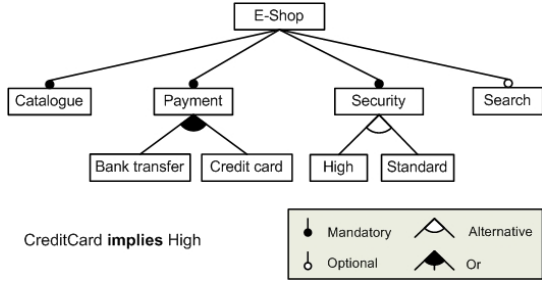


Figure 1: A feature model by Sergio Segura, Wikipedia.

For this paper, cross-tree constraints are formulas in propositional logic where features are the variables in the formulas. Two examples of common cross-tree constraints are $A \implies B$ (A requires B) and $A \implies \neg B$ (A excludes B).

2.2 SAT-based Analysis of Feature Models

The goal of SAT-based analysis of FMs is to find an assignment to the features such that the structural and cross-tree constraints are satisfied. It turns out that there is a natural reduction from feature models to SAT [2]. Each feature is mapped to a Boolean variable, the variable is true/false if the feature is selected/deselected. The structural and cross-tree constraints are encoded as propositional logic formulas. The SAT solver can answer questions like whether the feature model is void, when the model captures no products. The solver can also be adapted to handle product configuration: given a set of features that must be present and another set of features that must be absent, the solver will find a product that satisfies the request or answer that no such product exists. Optimization is also possible such as finding the product with the highest performance, although for this we need optimization solvers. Dead features, features that cannot exist in any valid products, can also be detected using solvers. More generally, SAT solvers provide a variety of possibilities for automated analysis of FMs, where manual analysis may be infeasible. Many specialized solvers [11, 12, 27] for FM analysis have been built that use SAT solvers as a backend. It is but natural to ask why SAT-based analysis tools scale so well and are so effective in diverse kinds of analysis of large real-world FMs. This question has been studied with randomly generated FMs based on realistic parameters [18, 24, 22] where all the instances were easily solved by a modern SAT solver. In this paper, we are studying large real-world FMs to explain why they are easy for SAT solvers.

3. EXPERIMENTS AND RESULTS

In this section, we describe the experiments that we conducted to better understand the effectiveness of SAT-based analyses of FMs. We assume the reader is familiar with the translation of FMs to Boolean formulas in conjunctive normal form (CNF).

3.1 Experimental Setup and Benchmarks

All the experiments were performed on 3 different comparable systems whose specs are as follows: Linux 64 bit machines with 2.8 GHz processors and 32 GB of RAM.




Figure 2: Solving times for real-world feature models with Sat4j.

Table 1 lists 15 real-world feature models translated to CNF from a paper by Berger et al. [7]. The number of variables in these models range from 544 up to 62470, and the number of clauses range from 1020 to 343944 clauses. Three of the models, the ones named 2.6.*, represent the x86 architecture of the Linux kernel. A clause is binary if it contains exactly 2 literals. If every clause is binary, then the satisfiability problem is called 2-SAT and it is solvable in polynomial time. A clause is Horn/anti-Horn if it contains at most one positive/negative literal. If every clause is Horn, then the satisfiability problem is called Horn-satisfiability and it is also solvable in polynomial time (likewise for anti-Horn). Binary/Horn/anti-Horn clauses account for many of the clauses, but not overwhelmingly⁴. Given the size of these real-world FMs, it is not a priori obvious that they are easy for modern CDCL solvers to solve.

3.2 Experiment: How easy are real-world FMs

Figure 2 shows the solving times of the real-world feature models with the Sat4j solver [15] (version 2.3.4). The implementation is known to be on the slower end of SAT solvers. With the default settings, the hardest feature model took a little over a second to solve. Clearly, the size of these feature models is not a problem for Sat4j. Modern SAT solvers have (at least) 4 main features that contribute to its performance: conflict-driven clause learning with backjumping, random search restarts, Boolean constraint propagation (BCP) using lazy data structures, and conflict-based adaptive branching called VSIDS [14]. Sat4j implements all these features. Figure 2 shows the running time after turning off 3 of these features, but the running times did not suffer significantly. BCP is surprisingly effective for real-world feature models. This result shows that the modern breakthroughs in SAT solving have very little impact for FMs. The efficiency even without state of the art techniques suggests that real-world feature models are easy to solve.

3.3 Experiment: Constructing hard artificial FMs

The question we posed in this experiment was whether it is possible to construct hard artificial FMs, and if so what would their structural features be. Indeed, we were able to

⁴Lots of Horn and anti-Horn clauses does not necessarily imply a problem is easy. For example, every clause in 3-SAT is either Horn or anti-Horn yet we currently do not know how to solve 3-SAT efficiently in general.

Model	Variables	Clauses	Horn (%)	Anti-Horn (%)	Binary (%)	Other (%)
2.6.28.6-icse11	6888	343944	7.54	50.80	6.19	47.29
2.6.32-2var	60072	268223	61.81	70.66	27.02	5.01
2.6.33.3-2var	62482	273799	63.24	70.05	27.74	5.08
axTLS	684	2155	71.04	52.99	25.85	7.19
buildroot	14910	45603	76.78	61.81	40.24	2.68
busybox-1.18.0	6796	17836	79.79	56.91	37.25	1.94
coreboot	12268	47091	82.67	68.53	45.73	0.59
ecos-icse11	1244	3146	92.59	73.36	73.27	6.64
embtoolkit	23516	180511	29.06	87.65	17.81	0.34
fiasco	1638	5228	84.87	52.85	38.49	1.42
freebsd-icse11	1396	62183	8.25	84.79	2.63	9.52
freetz	31012	102705	76.56	52.35	34.60	3.04
toybox	544	1020	90.49	67.75	46.76	0.00
uClinux-config	11254	31637	69.23	63.42	30.73	0.96
uClinux	1850	2468	100.00	75.04	50.00	0.00

Table 1: Clause and variable count of real-world FMs. The last four columns counts the percentage of clauses with the specified property. “Other” are clauses that are neither Horn, anti-Horn, nor binary. The percentages do not sum to 100% because Horn, anti-Horn, binary, and other are not disjoint.

construct small feature models that are very hard for SAT solvers. The procedure we used to generate such models is as follows:

1. First, we randomly generate a small and hard CNF formula. One such method is to generate random 3-SAT with a clause density of 4.25. It is well-known that such random instances are hard for SAT solver to solve. We then used such generated problems as the cross-tree constraints for our hard FMs.
2. Second, we generate a small tree with only optional features. The variables that occur in the cross-tree constraints generated in the first step are the leaves of the tree.

The key idea in generating such FMs is that for any pair of variables in the cross-tree constraints, the tree does not impose any constraints between the two. The problem then is as hard as the cross-tree constraints because a solution for the feature model is a solution for the cross-tree constraints and vice-versa. Using this technique, we can create small feature models that are hard to solve. Unlike these hard artificial FMs, the cross-tree constraints in large real-world FMs are evidently easy to solve.

3.4 Experiment: Variability in real-world FMs

Feature models capture variability, and we believe that the search for valid product configuration is easier as variability increases. We hypothesized that finding a solution is easy, when the solver has numerous solutions to choose from. We ran the feature models with sharpSAT [28], a tool for counting the exact number of solutions. The results are in Table 2. We found that real-world FMs display very high variability, i.e., have lots of solutions. For comparison, the 9 BMC models in the sharpSAT evaluation [28] have roughly the same number of variables and the same number of solutions as the FMs. The average solving times of the BMC models is 1.19 seconds with Sat4j.

Model	Variables	Solutions
axTLS	684	4.2×10^{20}
busybox-1.18.0	6796	8.4×10^{216}
coreboot	12268	1.4×10^{94}
ecos-icse11	1244	4.9×10^{125}
embtoolkit	23516	2.1×10^{218}
fiasco	1638	3.5×10^{14}
freebsd-icse11	1396	8.3×10^{313}
toybox	544	1.4×10^{17}
uClinux-config	11254	7.7×10^{417}
uClinux	1850	1.6×10^{91}

Table 2: The variability in real-world feature models, counted by sharpSAT. The computation finished for 10 of the 15 models.

Variability is the reason feature models exist. The feature groups and optional features increase the variability in the model, and results above suggests the variability grows exponentially with the size of the model measured in number of variables. High variability suggests that real-world FMs should have structural properties that make them easy to solve.

3.5 Experiment: Why solvers perform very few backtracks for real-world FMs

The goal of this experiment was to ascertain why solvers make so few backtracks while solving real-world FMs. First observe that when a solver needs to make new decision, it must guess the correct assignment for the decision variable-in-question, under the current partial assignment, in order to avoid backtracking. Also observe that if a decision variable is unrestricted (defined above) then either a *true* or a *false* assignment to such a variable would lead to a satisfying assignment. In other words, a preponderance of unrestricted variable in an input formula to a solver implies that the solver will likely make few mistakes in assigning values to decision variables and thus perform very few backtracks during solving.

Given the above line of reasoning, we designed an experiment that would increase our confidence in our hypothesis that a vast majority of the variables in real-world FMs are unrestricted. The experiment is as follows: whenever the solver branches on a decision or backtracks (i.e., when the partial assignment changes), we take a snapshot of the solver’s state. We want to examine how many unassigned variables are unrestricted/restricted. This requires copying the state of the snapshot into a new solver and checking if there are indeed solutions in both branches of the variable under the current partial assignment, in which case it is unrestricted. If only one branch has a solution, then it is restricted. If neither branch has a solution, then the current partial assignment is unsatisfiable.

Figure 3 shows how the unrestricted variables for one Linux-based feature model change over the course of the search. Time is measured by the number of decisions and conflicts. The y-axis shows the unassigned variables under the current partial assignment. The red area denotes unrestricted variables. If the solver branches on a variable in the red region, the solver will remain on the right track to finding a solution. The blue/green area denotes restricted variables. If the solver branches on a variable in the blue/green region, the solver must assign the variable the correct value⁵. The pink area means the current partial assignment is unsatisfiable and the solver will need to backtrack eventually. The 15 real-world feature models have very large red regions.

A large amount of unrestricted variables suggests that the instance is easy to solve. When the decision heuristic branches on an unrestricted variable it does not matter which branch to take, the partial assignment will remain satisfiable either way. Feature models offer enough flexibility such that the SAT solvers rarely run into dead ends.

3.6 Experiment: Simplifications

We hypothesized that since the vast majority of the variables are unrestricted they can be easily simplified away. Furthermore, the remaining variables/clauses, that we call a “core” would be small enough such that even a brute-force approach could solve it easily.

We implemented a number of standard simplifications that are particularly suited for eliminating variables. These simplifications were implemented as a preprocessing step to the Sat4j solver. Briefly, a simplification is a transformation on the SAT instance where the output is smaller and equisatisfiable to the input. We carefully chose simplifications that run in time polynomial in the number of variables in the input formula.

What we found was that indeed the simplifications were effective in eliminating more than 99% of the variables from real-world FMs. In 11 out of 15 instances from our benchmarks, the simplifications completely solved the instance without resorting to calling the backend solver. In the remaining cases, the cores were very small (atmost 53 vari-

⁵If the solver picks a random assignment for a restricted variable, then the solver will make a correct choice 50% of the time. In practice, SAT solvers bias towards false so negatively restricted variables might be preferable to positively restricted variables. The bias is often configurable.

ables) and were largely Horn clauses that were easily solved by Sat4j.

The simplifications we implemented are described below:

3.6.1 Equivalent Variable Substitution

DEFINITION 3.6.1. *If $x \implies y$ and $y \implies x$ where x and y are variables, then replace x and y with a fresh variable z .*

The idea behind this technique is to coalesce the variables since effectively $x = y$. This simplification step is useful for mandatory features that produces bidirectional implication in the CNF translation. This step reduces the number of variables by one and reduces the length of clauses containing x and y by one. Equivalent literal substitution is a more powerful variant, but the results will show that equivalent variable substitution is enough.

3.6.2 Subsumption

DEFINITION 3.6.2. *If $C_1 \subset C_2$ where C_1 and C_2 are clauses, then C_2 is subsumed by C_1 . Remove all subsumed clauses.*

The idea here is that subsumed clauses are trivially redundant. Initially, no clauses are subsumed in the real-world feature models. After other simplifications, some clauses will become subsumed.

3.6.3 Self-Subsuming Resolution

DEFINITION 3.6.3. *If*

$$\begin{aligned} C \vee x \\ C \vee \neg x \vee D \end{aligned}$$

where C and D are clauses and x is a literal, then replace with

$$\begin{aligned} C \vee x \\ C \vee D \end{aligned}$$

The idea here is that if $x = \text{true}$ then $C \vee \neg x \vee D$ reduces to $C \vee D$. If $x = \text{false}$ then $C = \text{true}$ in which case both $C \vee \neg x \vee D$ and $C \vee D$ are satisfied. In either case, $C \vee \neg x \vee D$ is equal to $C \vee D$, hence the clause can be shortened by removing the variable x .

Some features in the 3 Linux models are tristate: include the feature compiled statically, include the feature as a dynamically loadable module, or exclude the feature. Tristate features require two Boolean variables to encode. This simplification step is particularly useful for tristate feature models. For example, the feature A is encoded using Boolean variables a and a' . The table below shows how to interpret the values of this pair of variables.

a	a'	Meaning
<i>true</i>	<i>true</i>	Include A compiled statically
<i>true</i>	<i>false</i>	Include A as a dynamically loadable module
<i>false</i>	<i>false</i>	Exclude feature A

Table 3: Interpretation of tristate features.

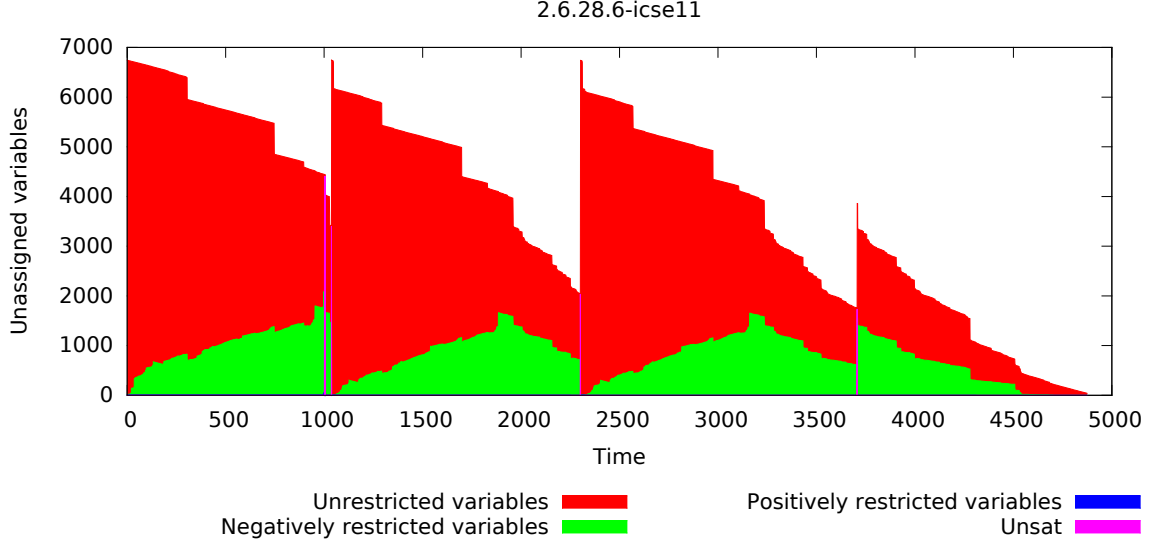


Figure 3: Unrestricted and restricted variables for a Linux-based feature model.

To restrict the possibilities to one of these 3 combinations, the translation adds the clause $a \vee \neg a'$. Since the interpretation of the feature requires both variables, they often appear together in clauses. Any other clause that contains $a \vee \neg a'$ will be removed by subsumption. Any other clause that contains $a \vee a'$ or $\neg a \vee \neg a'$ will be shortened by self-subsuming resolution.

3.6.4 Variable Elimination

DEFINITION 3.6.4. Let T be the set of all clauses that contain a variable and its negation. Let x be a variable.

S_x is the set of clauses where the variable x appears only positively. $S_{\bar{x}}$ is the set of clauses where the variable x appears only negatively. The variable x is eliminated by replacing the clauses S_x and $S_{\bar{x}}$ with

$$\{C_1 \vee C_2 \mid (x \vee C_1) \in S_x, (\neg x \vee C_2) \in S_{\bar{x}}, (C_1 \vee C_2) \notin T\}$$

The idea here is to proactively apply the *resolution* rule, a well known inference rule from propositional logic. As a result, the variable x no longer appears in the formula. A variable is only eliminated if the number of output clauses is less than or equal to the number of input clauses.

This simplification step is very effective for pure literals. A variable is pure if it appears either only positively or only negatively in the input formula. If a variable is pure, then variable elimination will eliminate that variable and every clause containing that variable. 37.8% of variables in the 15 real-world feature models from Table 1 are pure. More variables can become pure as the formula is simplified.

3.6.5 Asymmetric Branching

DEFINITION 3.6.5. For a clause $x \vee C$, where x is a literal and C is the remainder of the clause, temporarily add the constraint $\neg C$. If a call to BCP returns unsatisfiable,

then learn the clause C . The new learnt clause subsumes the original clause $x \vee C$ so remove the original clause. Otherwise, no changes.

3.6.6 RCheck

DEFINITION 3.6.6. For a clause C , temporarily replace the clause with the constraint $\neg C$. If a call to BCP returns unsatisfiable, then the other clauses imply C so the clause is redundant. Remove C from the formula. Otherwise, no changes.

The idea is to remove all clauses that are implied modulo BCP. Implied clauses can be useful for a SAT solver to prune the search space, for example learnt clauses, but they complicate analysis.

3.6.7 BCP

The last simplification step is to apply BCP. Some of the other steps shorten clauses, so take advantage of any newly discovered unit clauses

3.6.8 Fixed Point

We repeat the simplifications a maximum of 5 times or stop when a fixed point is reached. The simplifications can only shorten the size of the input formula with the exception of variable elimination. It is unclear how many passes of simplification is necessary to reach a fixed point in the worst-case with variable elimination in the mix. The upper limit of 5 passes is to guarantee a polynomial (in the size of the input formula) number of passes. For the models from Table 1, only 2 to 3 passes are needed to reach a fixed point using the current implementation built on top of Minisat's simplification routine. Each additional pass, in practice, experiences diminishing returns and 5 passes should be sufficient to simplify real-world feature models. The final output is the core.

Model	Variables	Clauses	Horn (%)	Anti-Horn (%)	Binary (%)	Other (%)
2.6.28.6-icse11	30	335	93.73	6.87	93.73	0.00
2.6.32-2var	0	0	NA	NA	NA	NA
2.6.33.3-2var	0	0	NA	NA	NA	NA
axTLS	0	0	NA	NA	NA	NA
buildroot	0	0	NA	NA	NA	NA
busybox-1.18.0	0	0	NA	NA	NA	NA
coreboot	53	153	84.31	73.86	58.17	0.00
ecos-icse11	0	0	NA	NA	NA	NA
embtoolkit	20	67	82.09	22.39	17.91	0.00
fiasco	30	384	99.74	7.03	99.74	0.00
freebsd-icse11	0	0	NA	NA	NA	NA
freetz	0	0	NA	NA	NA	NA
toybox	0	0	NA	NA	NA	NA
uClinux-config	0	0	NA	NA	NA	NA
uClinux	0	0	NA	NA	NA	NA

Table 4: Clause and variable count of simplified real-world FMs.

3.6.9 Time

The worst-case execution time of the simplifications is polynomial in the size of the input formula. The time is measured in regards to the size of the Boolean satisfiability input formula. To prove that the time is also polynomial in the size of the feature model, then the size of CNF must be polynomial in the size of the feature model. The parent-child relationship of mandatory/optional features translates to a linear number of binary clauses. Non-default feature groups, encoded as a Boolean cardinality constraint, introduces a polynomial (in the size of the feature group) number variables and clauses [1]. The cross-tree constraints are written in propositional logic and the translation to CNF increases the size of the formula linearly via Tseitin transformation. Therefore, the size of the translated CNF is at most polynomial in size of the input feature model.

3.6.10 Simplified Feature Models

Table 4 shows the feature models after simplification. 11 of the models are solved completely. The remaining 4 are small in terms of variables and clauses. The feature model “coreboot” is the largest in terms of variables.

At least 80% of the clauses in the simplified models are Horn. Horn-satisfiability can be solved in polynomial time. The algorithm works by applying BCP, and the Horn-formula is unsatisfiable if and only if the empty clause is derived. A similar algorithm for anti-Horn-satisfiability also exists. BCP is the engine for solving Horn-satisfiability in polynomial time and modern SAT solvers implement BCP, hence we hypothesize that Boolean formulas with ≤ 53 variables and $\geq 80\%$ Horn clauses are easy for SAT solvers to solve.

Figure 5 shows how Horn clauses make solving easier. Note that for 3-SAT, every clause is either Horn xor Antihorn, hence the symmetry in the graph. Every point in the graph is the average running time of 100 randomly generated 3-SAT instances. The instances were generated with 200 variables, which is larger than the simplified “coreboot” model. The running times where the Horn clauses exceed 80% is very low. Clause learning, adaptive branching, and random restarts are not necessary for this case. Although we originally suspected the core might be hard for its size, the core

itself turned out to be easy as well.

3.6.11 Application

The above simplifications have high overhead. If the model is only solved once, then simplifying first will increase the running time anywhere between 3 to 30 fold. However, some types of analyses solve the feature models many times. For these circumstances, it might be beneficial to simplify the original model once so the overhead is spread over the runs. Unfortunately, these types of analyses will add additional clauses after each run, in which case eliminated variables might be reintroduced and variable elimination is unsound. The variable elimination step will need redo part of its simplification. An alternative is to avoid variable elimination and equivalent variable substitution. Unfortunately, variable elimination is by far the most effective simplification (see Figure 4) and avoiding it will greatly impact the effectiveness of simplifications.

3.7 Experiment: Treewidth of real-world FMs

Experiments by Pohl et al. [22] show treewidth of the CNF representation of randomly-generated FMs to be strongly correlated with their corresponding solving times. We repeat the experiment on the real-world FMs to see if the correlation exists. In our experiments, treewidth is computed by finding a lower and upper bound because the exact treewidth computation is too expensive to compute. The longer the calculation runs, the tighter the bounds are. We used the same algorithm and package used by Pohl et al. We gave the algorithm a timeout of 3600 seconds and 24 GB of heap memory, up from 1000 seconds and 12 GB in the original experiment by Pohl et al. Figure 6 shows the results. The computation failed to place any upper bound on 9 FMs, and we omit these results because we do not know how close the computed lower bounds are to the exact answer. For the 6 models in the table, the lower and upper bound are close, and hence close to the exact answer.

We used Spearman’s rank correlation, the same correlation method as in the original experiment, to correlate the lower bound and time. We found the correlation between the lower bound and time to be 0.2571429. We get the same number when we replace the lower bound with the upper




Figure 4: The effectiveness of simplifications at reducing the number of variables. We try turning off simplifications one at a time and measure the result. The line “simplifications without variable elimination” is high on the y-axis, hence simplifying is much less effective at reducing the number of variables without variable elimination.




Figure 5: Solving Horn clauses for randomized 3-SAT (200 variables and 850 clauses) with Sat4j. The clause density is set to 4.25, where random 3-SAT is hardest to solve.

bound (the method is based on the rank, the lower and upper bound rank the models in the same order). Our results show a significantly poorer correlation than previous work on randomly-generated FMs. Although our sample is small, research on treewidth in real-world SAT instances (not FMs) have also found similar results, i.e., that treewidth of input Boolean formulas is not strongly correlated with running time of solvers and is not indicative of an instance’s hardness [17].

4. THREATS TO VALIDITY

In this Section we address threats to validity of our experimental methodology and results.

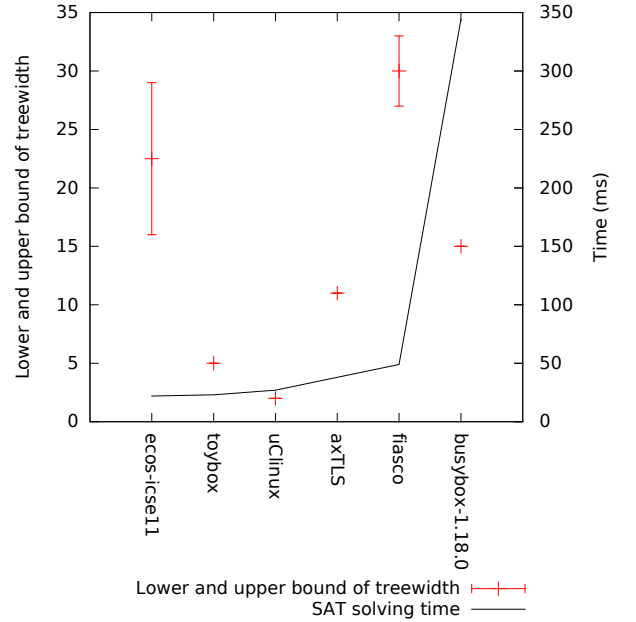


Figure 6: Treewidth of real-world FMs plotted against solving time. 4 models have exact bounds.

- **Validity of FMs used for the Experiments:** The studied collection of large real-world feature models [7] may not be representative of all large real-world feature models. This is indeed a concern since we focused primarily on software product lines, and there is no *a priori* reason to believe that the characteristics true of these FMs carry over to other FMs for other kinds of product lines, such as for cars. However, the collection includes all large real-world models that are publicly available. Further, systems software, the domain of these models, is known to produce some of the most complex configuration constraints [5]. Hence *a posteriori*, after taking into account the complexity of these models, we believe that these findings will likely hold for many other large real-world models.
- **Validity of Experimental Methodology:** The unrestricted variables experiment takes the partial assignments we encounter during the run of the solver with default parameters. The partial assignments could vary if the branching heuristic initiated with a random seed. Ideally, we would like to show that unrestricted variables are plentiful under any partial assignment but this would require an infeasible amount of effort to prove.
- **Validity of Sample Size:** A potential weakness of our experiments is the sample size. We used all the publicly available benchmarks in our experiments. Unfortunately, there are only 15 of them. Furthermore, for the treewidth experiments, the models were so big that only 6 of the 15 completed. Even though we only managed to obtain a small sample, we believe that our experiments are valid for the vast majority of the real-world FMs. Our reasoning is as follows: FMs by their very nature have high variability (borne out in our experiments as well). This variability reflects in the structural characteristics of FMs, such as shallow trees with lots of loosely-constrained optional features as leaves. These optional features in turn correspond to unrestricted variables in the respective Boolean representation of FMs, which are the cause for the “easiness” of solving these models. These conclusions follow largely from our understanding of semantics of FMs. Our experiments provide further evidence, and bolster these conclusions.
- **Correctness of Mapping of FMs into SAT:** The analyses rely on the correctness of mapping of the Kconfig and CDL models to propositional logic. We rely on existing work that reverse engineered the semantics of these languages from documentation and configuration tools and specified the semantics formally [8, 7].
- **Validity of Translation Strategies:** The results could be influenced by the chosen translation strategies to translate the FMs into SAT formulas. However, these translation strategies are standard and used in innumerable other domains where problems are translated into SAT formulas, to be then solved by SAT solvers.
- **Validity of Choice of Solvers:** All our experiments were performed using the Sat4j solver [15]. An important concern is whether these results would apply to other SAT solvers and analysis tools built using them. Our choice of Sat4j was driven by the fact that it is one

of the simplest conflict-driven clause-learning (CDCL) SAT solvers publicly available, and most other SAT solvers used in FM analysis are conflict-driven clause-learning. Hence, we believe that our results would apply equally well to other solvers used in FM analysis.

- **Comparison with CSP and BDD-based Solvers:** We plan to extend our work by comparing the performance of SAT solvers for FM analysis over large real-world models, with that of CSP and BDD-based solvers. In this work, we focused only on SAT solvers as it became evident that the heuristics employed by them (e.g., resolution) are very effective in exploiting the most salient features of large real-world FMs.

5. RELATED WORK

Using constraints solvers to analyze feature models has a long tradition. Benavides et al. [3] give a comprehensive survey of techniques for analyzing feature models, such as checking consistency and detecting dead features, and for supporting configuration, such as propagating feature selections. These techniques use a range of solvers, including SAT solvers, Constraint Satisfaction Problem (CSP) solvers, and Binary-Decision Diagram (BDD) packages. Batory [2] was first to suggest the use of SAT solvers to support feature model analyses.

Several authors have investigated the efficiency of different kinds of solvers for feature model analyses. Benavides et al. [4] compare the performance of SAT, CSP, and BDD solvers on two sample analyses on randomly generate feature models with up to 300 features. They show that BDD solvers are prone to exponential growth in memory usage on larger models, while SAT solvers achieve good runtime performance in the experiments. Pohl et al. [21] run similar comparison on feature models from the SPLOT collection, including a wider set of solver implementations. SPLOT models are relatively small—the largest one has less than 300 features—and are derived from academic works. Their results show that SAT solvers work well also for SPLOT models, although they detect some performance variation for C- vs. Java-based SAT solvers; they also confirm the tractability challenges for BDD solvers. Mendonca et al. [19] achieve scalability of BDD-based analyses to randomly generated feature models with up to 2000 features by optimizing the translation from the models to BDDs using variable ordering heuristics based on the feature hierarchy. Further, Mendonca et al. [18] show that SAT-based analyses scale to randomly generated feature models with up to 10000 features. All this previous work shows that SAT solvers perform well on small realistic feature models (SPLOT collection) and large randomly generated feature models. Although SAT solvers have been successfully applied to analyze the feature model of the Linux kernel (e.g., [26]), we are unaware of prior work systematically studying the performance of SAT solvers on large (with 1000 or more features) real-world feature models, which is what we focus on in this paper.

Previous work has also investigated the hardness of the SAT instances derived from feature models, aiming at more general insights into the tractability of SAT-based analyses. Mendonca et al. [18] have studied the SAT solving behavior of randomly generated 3-SAT feature models. A 3-SAT

feature model is one whose cross-tree constraints are 3-SAT formulas. While random 3-SAT formulas become hard when their clause density approaches 4.25 (so-called phase transition), this work shows experimentally that 3-SAT feature models remain easy across all clause densities of their corresponding 3-SAT formulas. Our work is different since it investigates the hardness of large real-world feature models rather than random ones. Even though random feature models are easy on average, specific instances may still be hard; for example, a tree of optional features with a hard 3-SAT cross-tree formula over the tree leaves is, albeit contrived, a hard feature model. Segura et al. [24] treat finding hard feature models as an optimization problem. They apply evolutionary algorithms to generate feature models of a given size that maximize solving time or memory use. Their experiments show that relatively small feature models can become intractable in terms of memory use for BDDs; however, the approach did not generate feature models that would be hard for SAT solvers. Again, this work only considers synthetically generated feature models, rather than real-world ones.

Finally, Pohl et al. [22] propose graph width measures, such as treewidth, applied to graph representations of the CNF derived from a feature model as an upper bound of the complexity of analyses on the model. Their work evaluates this idea on a set of randomly generated feature models with up to 1000 features using SAT, CSP, and BDD solvers; they also repeat the experiment on the SPLOT collection. They find a significant correlation between certain treewidth measures on the incidence graph and the running time for most of the solvers. Further, the authors note that it is still unclear why SPLOT models are easier than generated ones, and whether that observation would hold for large-scale real-world feature models. Our work addresses this gap by investigating the hardness of large-scale real-world feature models and providing an explanation why they are easy. Moreover, we were not able to identify any correlation between the treewidth and easiness of the SAT instances derived from large real-world models, which calls for more work to find effective hardness measures for feature models.

Large real-world models have become available to researchers only recently. While some papers hint at the existence of very large models in industry [6], these models are typically highly confidential. Sincero [25] was first to observe that the definition of the Linux kernel build configuration, expressed in the Kconfig language, can be viewed as a feature model. Berger et al. [8] identify the Component Configuration Language (CDL) in eCos, an open-source real-time operating system, as additional feature modeling language and subsequently [7] create and analyze a collection of large real-world feature models from twelve open-source projects in the software systems domain. We use this collection as a basis for our work.

A related field is the application of SAT solvers to software package management and configuration. For example, Vouillon and Di Cosmo [29] apply SAT solvers to undesirable dependencies in Debian package structures and Le Berre and Rapicault [16] use SAT solvers to aid the package installation and upgrades for Eclipse. While package dependencies are related to feature dependencies in feature models (e.g.,

[10]), the general structure of package repositories and feature models are different: the former are graphs managed de-centrally, while the latter have centrally-managed hierarchies [5]. Yet, the SAT instances generated by both problems may share characteristics and the proposed simplifications may be applicable in both settings with comparable results; we leave this question for future work.

Researchers have also studied finding optimal configurations for a given feature model. For example, Sayyad et al. [23] use evolutionary algorithms to generate approximate solutions to the problem of finding optimal configurations. Such optimization problems are harder than consistency checking, but the simplification techniques presented in our paper may be useful for speeding up exact optimization using solvers (e.g., [20]).

6. CONCLUSIONS

In this paper, we provided an explanation, with strong experimental support, for the scalability of SAT solvers on large real-world FMs. The explanation is that the overwhelming majority of the variables in real-world FMs are *unrestricted*, and solvers tend not to backtrack in the presence of such variables. We argue that the reason for the presence of large number of unrestricted variables in real-world FMs has to do with high variability in such models. We also found that if we switch off all the heuristics in modern SAT solvers, except Boolean constant propagation (BCP) and backjumping (no clause-learning), then the solver does not suffer any deterioration in performance while solving FMs.

Moreover, we ran a set of simplifications with substantial reduction to the size of the instances. In fact, a majority of the models were solved outright. The most effective simplification, variable elimination, is based on binary resolution. While these simplification do incur an overhead, they can nonetheless be very useful if used judiciously in conjunction with a CDCL solver.

Visit this URL to download the tools and data:
<https://ece.uwaterloo.ca/~vganesh/featuremodelsat.html>.

7. REFERENCES

- [1] O. Bailleux and Y. Bouffekh. Efficient CNF encoding of boolean cardinality constraints. In F. Rossi, editor, *Principles and Practice of Constraint Programming - CP 2003*, volume 2833 of *Lecture Notes in Computer Science*, pages 108–122. Springer Berlin Heidelberg, 2003.
- [2] D. Batory. Feature models, grammars, and propositional formulas. In *Proceedings of the 9th International Conference on Software Product Lines, SPLC'05*, pages 7–20, Berlin, Heidelberg, 2005. Springer-Verlag.
- [3] D. Benavides, S. Segura, and A. Ruiz-Cortés. Automated analysis of feature models 20 years later: A literature review. *Information Systems*, 35(6):615 – 636, 2010.
- [4] D. Benavides, S. Segura, P. Trinidad, and A. Ruiz-Cortés. A first step towards a framework for the automated analysis of feature models. In

Managing Variability for Software Product Lines: Working With Variability Mechanisms, 2006.

- [5] T. Berger, R.-H. Pfeiffer, R. Tartler, S. Dienst, K. Czarnecki, A. Wasowski, and S. She. Variability mechanisms in software ecosystems. Under revision.
- [6] T. Berger, R. Rublack, D. Nair, J. M. Atlee, M. Becker, K. Czarnecki, and A. Wasowski. A survey of variability modeling in industrial practice. In *Proceedings of the Seventh International Workshop on Variability Modelling of Software-intensive Systems, VaMoS '13*, pages 7:1–7:8, New York, NY, USA, 2013. ACM.
- [7] T. Berger, S. She, R. Lotufo, A. Wasowski, and K. Czarnecki. A study of variability models and languages in the systems software domain. *IEEE Trans. Softw. Eng.*, 39(12):1611–1640, Dec. 2013.
- [8] T. Berger, S. She, R. Lotufo, A. Wasowski, and K. Czarnecki. Variability modeling in the real: A perspective from the operating systems domain. In *Proceedings of the IEEE/ACM International Conference on Automated Software Engineering, ASE '10*, pages 73–82, New York, NY, USA, 2010. ACM.
- [9] K. Czarnecki and U. W. Eisenecker. *Generative Programming: Methods, Tools, and Applications*. ACM Press/Addison-Wesley Publishing Co., New York, NY, USA, 2000.
- [10] R. Di Cosmo and S. Zacchiroli. Feature diagrams as package dependencies. In *Proceedings of the 14th International Conference on Software Product Lines: Going Beyond, SPLC'10*, pages 476–480, Berlin, Heidelberg, 2010. Springer-Verlag.
- [11] R. Gheyi, T. Massoni, and P. Borba. A theory for feature models in alloy. pages 71–80, Portland, USA, November 2006.
- [12] M. Janota. Do SAT solvers make good configurators? In *SPLC (2)*, pages 191–195, 2008.
- [13] K. C. Kang, S. G. Cohen, J. A. Hess, W. E. Novak, and A. S. Peterson. Feature-oriented domain analysis (FODA) feasibility study. Technical report, Carnegie-Mellon University Software Engineering Institute, November 1990.
- [14] H. Katebi, K. A. Sakallah, and J. a. P. Marques-Silva. Empirical study of the anatomy of modern SAT solvers. In *Proceedings of the 14th International Conference on Theory and Application of Satisfiability Testing, SAT'11*, pages 343–356, Berlin, Heidelberg, 2011. Springer-Verlag.
- [15] D. Le Berre, A. Parrain, M. Baron, J. Bourgeois, Y. Irrilo, F. Fontaine, F. Laihem, O. Roussel, and L. Sais. Sat4j-a satisfiability library for java, 2006.
- [16] D. Le Berre and P. Rapicault. Dependency management for the eclipse ecosystem: Eclipse p2, metadata and resolution. In *Proceedings of the 1st International Workshop on Open Component Ecosystems, IWOCE '09*, pages 21–30, New York, NY, USA, 2009. ACM.
- [17] R. Mateescu. Treewidth in industrial SAT benchmarks. Technical report, Microsoft Research, 2011.
- [18] M. Mendonca, A. Wasowski, and K. Czarnecki. SAT-based analysis of feature models is easy. In *Proceedings of the 13th International Software Product Line Conference, SPLC '09*, pages 231–240, Pittsburgh, PA, USA, 2009. Carnegie Mellon University.
- [19] M. Mendonca, A. Wasowski, K. Czarnecki, and D. Cowan. Efficient compilation techniques for large scale feature models. In *Proceedings of the 7th International Conference on Generative Programming and Component Engineering, GPCE '08*, pages 13–22, New York, NY, USA, 2008. ACM.
- [20] R. Olacchia, S. Stewart, K. Czarnecki, and D. Rayside. Modeling and multi-objective optimization of quality attributes in variability-rich software. In *International Workshop on Non- functional System Properties in Domain Specific Modeling Languages (NFPinDSML'12)*, Innsbruck, Austria, 10/2012 2012.
- [21] R. Pohl, K. Lauenroth, and K. Pohl. A performance comparison of contemporary algorithmic approaches for automated analysis operations on feature models. In *Proceedings of the 2011 26th IEEE/ACM International Conference on Automated Software Engineering, ASE '11*, pages 313–322, Washington, DC, USA, 2011. IEEE Computer Society.
- [22] R. Pohl, V. Stricker, and K. Pohl. Measuring the structural complexity of feature models. In *Automated Software Engineering (ASE), 2013 IEEE/ACM 28th International Conference on*, pages 454–464, Nov 2013.
- [23] A. S. Sayyad, T. Menzies, and H. Ammar. On the value of user preferences in search-based software engineering: A case study in software product lines. In *Proceedings of the 2013 International Conference on Software Engineering, ICSE '13*, pages 492–501, Piscataway, NJ, USA, 2013. IEEE Press.
- [24] S. Segura, J. A. Parejo, R. M. Hierons, D. Benavides, and A. R. Cortés. Automated generation of computationally hard feature models using evolutionary algorithms. *Expert Syst. Appl.*, 41(8):3975–3992, 2014.
- [25] J. Sincero, H. Schirmeier, W. Schröder-Preikschat, and O. Spinczyk. Is The Linux Kernel a Software Product Line? In F. van der Linden and B. Lundell, editors, *Proceedings of the International Workshop on Open Source Software and Product Lines (SPLC-OSSPL 2007)*, Kyoto, Japan, 2007.
- [26] R. Tartler, D. Lohmann, J. Sincero, and W. Schröder-Preikschat. Feature consistency in compile-time-configurable system software: Facing the linux 10,000 feature problem. In *Proceedings of the Sixth Conference on Computer Systems, EuroSys '11*, pages 47–60, New York, NY, USA, 2011. ACM.
- [27] T. Thum, D. Batory, and C. Kastner. Reasoning about edits to feature models. In *Proceedings of the 31st International Conference on Software Engineering, ICSE '09*, pages 254–264, Washington, DC, USA, 2009. IEEE Computer Society.
- [28] M. Thurley. sharpsat: Counting models with advanced component caching and implicit bcp. In *Proceedings of the 9th International Conference on Theory and Applications of Satisfiability Testing, SAT'06*, pages 424–429, Berlin, Heidelberg, 2006. Springer-Verlag.
- [29] J. Vouillon and R. Di Cosmo. Broken sets in software repository evolution. In *Proceedings of the 2013 International Conference on Software Engineering*,

ICSE '13, pages 412–421, Piscataway, NJ, USA, 2013.
IEEE Press.