Reasoning about Probabilistic Defense Mechanisms against Remote Attacks

Martín Ochoa*, Sebastian Banescu†, Cynthia Disenfeld‡, Gilles Barthe§, Vijay Ganesh¶

*Singapore University of Technology and Design
†Technische Universität München
‡University of Toronto
§ IMDEA Software
¶ University of Waterloo

Abstract—Despite numerous countermeasures proposed by practitioners and researchers, remote control-flow alteration of programs with memory-safety vulnerabilities continues to be a realistic threat. Guaranteeing that complex software is completely free of memory-safety vulnerabilities is extremely expensive. Probabilistic countermeasures that depend on random secret keys are interesting, because they are an inexpensive way to raise the bar for attackers who aim to exploit memory-safety vulnerabilities. Moreover, some countermeasures even support legacy systems. However, it is unclear how to quantify and compare the effectiveness of different probabilistic countermeasures or combinations of such countermeasures. In this paper we propose a methodology to rigorously derive security bounds for probabilistic countermeasures. We argue that by representing security notions in this setting as events in probabilistic games, similarly as done with cryptographic security definitions, concrete and asymptotic guarantees can be obtained against realistic attackers. These guarantees shed light on the effectiveness of single countermeasures and their composition and allow practitioners to more precisely gauge the risk of an attack.

1. Introduction

Memory-safety vulnerabilities, such as writing outside of boundaries of buffers of programs written in C/C++, represent a high security risk for applications facing a network or malicious input, since they can be exploited to execute arbitrary code on a victim's machine [1]. Although due to its severity, several countermeasures have been proposed in the last 15 years to mitigate the risk of exploitation, a new generation of exploits taking advantage of code re-use techniques [2], [3], [4] pose challenges to practitioners and researchers. Also, as the difficulty to exploit buffer overflows (BOF) increases, attackers resort to exploiting other memory-safety vulnerabilities, such as format string vulnerabilities, heap overflows and buffer over-reads (as in the case of the recent Heartbleed exploit [5]).

Security researchers have come up with many countermeasures (we interchangeably use the terms defense mechanisms and countermeasures) against memory-based exploitation, broadly classified as probabilistic and non-probabilistic. Examples of non-probabilistic countermeasures include access control protection to certain memory regions to avoid remote code injection on the stack [1]. This kind of access control is currently implemented in most operating systems and at the hardware level by most CPU vendors. Unfortunately, such countermeasures cannot prevent code re-use attacks such as Return-oriented Programming (ROP) [2]. In ROP attacks, instead of injecting malicious code, attackers find sequences of instructions in the victim's code (called *gadgets*) that can be executed in arbitrary order by injecting pointers to them on the stack, which essentially allows running arbitrary code without violating the access control policy on the stack. Other non-probabilistic countermeasures, such as Control-Flow Integrity (CFI) [6] propose to statically compute a valid control-flow graph of a program and enforce run-time monitors that detect violations to this policy. Although promising, this technique implies source-code or binary instrumentation, and can have high performance overheads [7]. Moreover ROP attacks that respect the original control-flow graph can be launched even in contexts where CFI may be enforced [8].

Another class of countermeasures are probabilistic [1], and can intuitively be classified under the "software diversity" or "moving-target" defense mechanisms. The core idea behind them is that there is an element of randomness (a *seed* or a *key*), that changes with every execution of a program or with every fresh boot of the host executing the program, depending on the implementation. This nondeterminism is used with the goal to slow-down an attacker: without any further side-channel, an attacker would have to guess the secret in order to launch the attack. Depending on the size of the secret, this can effectively increase an attacker's effort to successfully carry out an attack, and demotivate him/her. Prominently, address space layout randomization (ASLR) [9] and random stack canaries [10] have found their way into most modern operating systems. Other countermeasures such as instruction-set randomization (ISR) [11] and PointGuard [12] have been proposed in the literature as valuable alternatives, because of ease of implementation (easy or no binary instrumentation) and low performance overheads. These characteristics are crucial for practical applications, where legacy software running on resource constrained environments such as industrial-control

systems pose several implementation challenges.

However, the effectiveness (or the lack thereof) of probabilistic countermeasures has been discussed in the literature for various reasons. First, as stated before, clever exploitation techniques such as ROP [2], Blind ROP [3] and JITROP [4] invalidate certain assumptions on the attacker, made by some countermeasures, rendering them less effective. On the other hand, side-channels attacks [13] and low entropy of randomization in certain architectures [14] make them practically insecure.

In this work, we make the observation that the guarantees provided (or intended) by probabilistic countermeasures for memory-safety vulnerabilities resemble the ones of cryptographic algorithms, and as such, can be represented and reasoned upon using state-of-the-art cryptographic tools and techniques. In particular, game-based reasoning [15] has gained momentum in the past decade as a useful technique to formalize proof techniques involving probabilistic guarantees such as semantic security [16] among others. At the core of our contribution is a security property definition that abstracts away from the concrete defense mechanism and low-level details (such as various assembly languages, processor architectures etc.), and that can be instantiated for various countermeasures. In particular, our modeling allows for reasoning about composed countermeasures in a natural way. We presented some of the key ideas discussed here in an earlier paper on *attack-resistance of defense mechanisms* [17], but have considerably improved upon and expanded these ideas in this paper.

Problem Statement. In this work, we address the problem of rigorously deriving security guarantees for probabilistic countermeasures against remote exploitation attacks for resource constrained but otherwise arbitrary attackers. We are motivated to provide formal guarantees of efficacy for defense mechanisms because such guarantees are crucial in characterizing and understanding why certain defense mechanisms work as intended while others fail.

Contributions. Our contributions can be summarized as follows:

- We quantify the security of a probabilistic countermeasure as the probability of an event in a probabilistic game, similar to how it is done for the formal verification of cryptographic primitives via security games [15], [18].
- We compute bounds for popular probabilistic countermeasures from the literature, showing that our approach generalizes. Crucially, we can analyze compositions of different countermeasures.
- We discuss how the use of replicas, similar as in Secure Multi-Execution (SME) [19] and DieHard [20] can be used to close certain side-channels in the implementation of probabilistic countermeasures. Additionally, we show that this technique can conceptually make exploitation harder against many countermeasures even if the keys on which the countermeasure is based are leaked.

The rest of this paper is organized as follows. Section 2 recaps fundamentals of stack-smashing attacks, since our running examples will exploit BOFs, and fundamentals of game-based cryptographic proofs. Section 3 presents the central concepts and definitions, and Section 4 shows applications to various countermeasures. We show how to reason on composed countermeasures in Section 5. Section 6 discusses how to plug-in leakage due to side-channels into our reasoning, and how this could be prevented by means of replicas. We discuss related work in Section 7 and conclude in Section 8.

2. Background

In the following, we provide some background on memory-based exploits and probabilistic countermeasures. We focus on stack smashing attacks because of their popularity and criticality, and since they will serve as a basis for our discussion. Finally, we summarize the fundamental concepts behind code-based cryptographic proofs.

2.1. Stack smashing attacks

For the sake of clarity of presentation, in this work we assume that the only way the attacker can violate the control-flow integrity of a remotely accessible program is via stack smashing attacks, e.g. code injection and code re-use. Stack smashing attacks involve writing data input by the attacker, beyond the declared boundaries of a statically allocated buffer, i.e. overflowing the buffer. If the input of the attacker is long enough, then it will overwrite the control-flow information (i.e. frame pointer and return address) of the program which is also stored on the stack in the vicinity of the overflowed buffer. Once the currently executing function of the program is finished it will continue execution from the instruction indicated by the overwritten return address. If this return address points to a valid instruction sequence, then the control-flow integrity of the program has been violated, since the program will deviate from its intended behaviour. Otherwise, if the return address points to a sequence of bytes that does not represent a valid instruction according to the instruction set architecture (ISA) or it points to an inaccessible memory address, then the program crashes.

Code injection is a type of stack smashing attack, where the attacker inserts the machine code s/he wants to execute in the input that is passed to the vulnerable program. This input is also known as the *exploit payload* or *shellcode* and it also contains other data additional to the machine code to be executed, e.g. the value with which the return address is to be overwritten. As described in the previous paragraph, this exploit payload will be written to the vulnerable buffer and beyond its declared boundary. A crucial step of the attack is therefore to overwrite the value of the return address with the address where this injected code is stored on the stack.

Code re-use is similar to code injection, however, the attacker inserts a sequence of addresses to existing code in the process memory of the target program. Each of these addresses generally points to a short sequence of instructions

ending with a return or jump instruction, which are called *gadgets*. In between the addresses of the gadgets that the attacker inserts, there may also be data values which are used by the instructions of the gadgets (e.g. `pop eax` takes such a data value from the stack). When a gadget finishes execution it will either (1) execute the return instruction which loads the address of the following gadget from the stack and continues execution, or (2) execute the jump instruction to an address (provided by the attacker via the exploit payload), to the following gadget. Again a crucial step for the attacker is to carefully craft the exploit payload to overwrite the return address with the address of the first gadget that s/he wants to execute.

2.2. System interface

In the previous subsection we have mentioned that the attacker can only remotely access the system on which the vulnerable program is running. To be more precise, we assume a synchronous communication channel such as TCP, i.e. after the attacker sends an input (request), s/he receives / observes an output (response) from the system. The output can be either a message sent directly by the vulnerable program or it can be an error message sent by the remote system, (e.g. time-out), due to a crash of the vulnerable program.

2.3. Code-based cryptographic proofs

Modern cryptography advocates the use of rigorous security definitions and reductionist proofs [16]. In a typical case, security definitions are captured by probabilistic experiments in which a challenger $\mathcal{C}$ interacts with an adversary $\mathcal{A}$, and the adversary will *win* the game if his response to the challenge computed by $\mathcal{C}$ fulfills a well-defined winning condition. An important class of experiments are the so-called indistinguishability games, where an adversary must guess some bit b sampled by the challenger; in this case, the winning condition is $b = b'$, where b' is the adversary's guess. Indistinguishability games can also be expressed by means of two experiments: in the first experiment, the challenger uses an ideal functionality, whereas in the second experiment, the challenger uses a real functionality. In both cases, the adversary returns a bit $b \in \{0, 1\}$. The advantage of the adversary is then defined as the distance (i.e. absolute value of the difference) between the probabilities that he returns 1 in the first experiment and in the second experiments. To manage the complexity of reductionist proofs, some cryptographers adopt the game-hopping approach, in which bounds on the adversary's advantage are proved by means of a sequence of transitions [15]. A further step is to write experiments using an imperative probabilistic programming language [15]. One can then give a formal semantics to probabilistic programs and use this semantics as a sound basis for rigorous mathematical reasoning [18]. One central tool in the game-hopping approach is reasoning up to failure event. In its simplest form, it considers two programs that are syntactically identical until a boolean

flag bad is set (a flag is a boolean variable that remains true once set).

Definition 1 (Identical until bad is set). Two games G_1 and G_2 are identical until bad is set, if both contain a bad flag (that is initially set to false) and they are *syntactically* identical but for a statement if $bad = true$ then S in one game and if $bad = true$ then T in the other, for some sequence of instructions T and S.

Equipped with the notion of *identical until bad is set*, we can proceed to state the Fundamental Lemma of Game Playing, which is a cornerstone of the hop-based technique proposed in [15], [18].

Theorem 1 (Fundamental Lemma of Game playing). Let G_1 and G_2 be two (terminating) games identical until bad is set. Let E an event defined on both games and let F denote the event $bad = true$. Then:

$$|\Pr[E]_{G_1} - \Pr[E]_{G_2}| \leq \Pr[F]_{G_i} \quad i = 1, 2$$

For a proof see [15].

Example 1 (PRF/PRP Switching Lemma). A simple example that illustrates the game playing principle is the Pseudo-Random Function/Pseudo-Random Permutation (PRF/PRP) switching lemma [15]. Let $Perm(n)$ be the set of all permutations over strings of n bits $\{0, 1\}^n$, and let $Func(n)$ the set of all functions from $\{0, 1\}^n$ to $\{0, 1\}^n$. The switching lemma states that the probability of an adversary distinguishing sampling from the two sets above given q oracle access calls to the sampling procedure is at most $\frac{q(q-1)}{2^{n+1}}$. Formally, let A^π be an adversary querying an oracle that samples a random permutation $\pi \in Perm(n)$, and A^ρ an adversary that queries an oracle sampling a random function $\rho \in Func(n)$. The adversary A^P, for $P \in \{\rho, \pi\}$ outputs a bit (0,1) after querying the random function/permutation q times. Then:

$$|\Pr[A^\pi = 1] - \Pr[A^\rho = 1]| \leq \frac{q(q-1)}{2^{n+1}}$$

We can model the two different oracles using the *pWhile* language of [18]. Here $x \xleftarrow{\$} S$ is a uniform random assignment from the set S. We assume that the set I is empty before giving oracle access to the adversary.

```
Proc. π
y ⟵$ {0,1}^n
if y ∈ I then
    y ⟵$ {0,1}^n \ I
I ← I :: y
return y
```

```
Proc. ρ
y ⟵$ {0,1}^n
return y
```

where $I :: y$ stands for concatenation of y to a list I. Now, we can syntactically transform these programs into the following, semantically equivalent ones (S_1 has the same output distribution as π and S_2 has the same output distribution as ρ):

```
Proc. S_1
y ←$ {0,1}^n
if y ∈ I then
    bad ← true
    y ←$ {0,1}^n \ I
I ← I :: y
return y
```
```
Proc. S_2
y ←$ {0,1}^n
if y ∈ I then
    bad ← true
I ← I :: y
return y
```

Since S_1 is identical to S_2 up to setting the bad flag to $true$, we can use the fundamental lemma of game playing and derive:

$$| \Pr[A^\pi = 1] - \Pr[A^\rho = 1]| \leq \frac{q(q-1)}{2^{n+1}}$$

$$| \Pr[A^\pi = 1] - \Pr[A^\rho = 1]| =$$
$$| \Pr[A^{S_1} = 1] - \Pr[A^{S_2} = 1]| \leq \Pr[bad = true]_{S_1}$$
$$\leq \frac{q(q-1)}{2^{n+1}}$$

since the probability of $y \in I$ is the probability of a collision in sampling q random elements, which is a standard birthday paradox bound.

Interpretation of bounds. Probabilistic experiments are always parametrized by a security parameter n; for instance, the security parameter can be related to the size of the underlying domain (in our example, the length of the bitstrings). Thus, the advantage of the adversary is implicitly given in terms of a function $\epsilon(n)$, and some care is required for interpreting the results—especially when the advantage is expressed relative to a security assumption. One standard interpretation is asymptotic: in this setting, one requires that adversaries are probabilistic polynomial time (PPT) algorithms and that their advantage is negligible in the security parameter, i.e. goes to zero exponentially fast. Another, and more appropriate for our purposes, interpretation is concrete security; in this case, one analyzes the values of $\epsilon(n)$ for chosen values of n that reflect practical scenarios.

3. Security of probabilistic defenses as games

In this section we present our formal approach on how to characterize the expected security guarantees for probabilistic defense mechanisms against memory-safety vulnerabilities and how to formalize them as code-based games. In the following we assume that the execution of the programs will eventually timeout and therefore they are terminating. Also we assume that programs are deterministic (by assuming their source of randomness is isolated as a program input). Thus, in this setting, programs can be though of as total functions.

3.1. Security definition

Let P be a low-level program for an x86-like architecture (i.e. x86 assembly). Let $[P] : \{0,1\}^* \to \{0,1\}^*$ be a semantics function for P. Let $[\![P]\!]$ be an ideal execution that will raise an exception and crash with a distinct output

crash whenever an attempt to exploit a memory-safety violation is detected at runtime. Formally there exists a set $\Omega(P) \in 2^{\{0,1\}^*}$ such that:

$$\forall \, \omega \in \Omega, \; [P](\omega) = \text{crash}$$

Let $\mathcal{A}$ be a PPT adversary that attacks P, which is potentially a vulnerable program. $\mathcal{A}$ knows the program P and can interact remotely with it (send inputs and can observe the respective outputs) by means of a machine executing it as discussed in Section 2.

Definition 2 (Effectiveness of defense mechanism). We say that a probabilistic countermeasure c is effective against remote exploitation if the probability of attacker $\mathcal{A}$ of distinguishing the execution of P with countermeasure c (denoted $[P + c]$), from the ideal execution $[\![P]\!]$ is negligible on a security parameter n. The adversary has oracle access to an instance of $[P + c]$, which he can query q times (where q is polynomial on n). Formally:

$$| \Pr[\mathcal{A}^{[P+c]} = i]| \leq \epsilon(n)$$

Where i is an input to P such that $i \in \Omega(P)$ and $[P + c](i) \neq \text{crash}$, for a negligible function ϵ.

Note that in this definition we abstract away from the concrete attack (i.e. executing a remote shell, or disabling stack execution protection), the vulnerability exploited (heap or stack overflow, use after free etc.) and the countermeasure used. We rely on the fact that a precondition for any attack is that the memory-safety violation goes undetected.

3.2. Game-based modeling

In order to instantiate Definition 2, we will write the core of a defense countermeasure as a probabilistic program (or game), as well as the ideal execution $[\![\cdot]\!]$.

The ideal program execution is:

```
Proc. [[P]](i)
if i ∈ Ω(P) then
    o ← crash
else
    o ← [P](i)
return o
```

where, for the sake of this paper, $\Omega(P)$ is the set of all inputs that will cause the program *to write out of bounds of a static buffer*. In principle our abstraction can be applied for other memory safety violations (such as heap overflows etc.) but we leave a thorough exploration of such extensions for future work.

Clearly

$$\Pr[i \in \Omega(P) \wedge o \neq \text{crash}] = 0$$

in this case.

Instead, a regular execution of a program P with buffer overflow vulnerabilities (together with an empty countermeasure $\emptyset$) can be abstracted as:

```
Proc. [P + ∅](i)
If i ∈ Ω(P) then
    ra ← i.payload[0]
    If ra ∈ Valid then
        o ← [M(P, ra)]
    else
        o ← crash
else
    o ← [P](i)
return o
```

where Valid is the set of valid return addresses. $\mathcal{M}(P, ra)$ is the sequence of bytes in the process memory of P, starting from address ra. Since this input would make the ideal execution crash, it means that it writes outside of bounds of a memory buffer. We assume that a function $i.\text{payload}[0]$ is given that extracts the part of the input that overwrites the return address in the stack. For simplicity we assume that this function always returns a value, that is, the adversary aims at overwriting the return address to perform some sort of code injection/reuse. Certainly, in particular cases the adversary might want just to manipulate local variables in the proximity of the vulnerable buffer, but in this work we limit ourselves to analyzing the more common and usually more critical attack that involves remote code execution.

Note that if the attacker overwrites the return address with a valid address, the exact output in this case is determined by the concrete state of the memory in the victim's host machine and could potentially result in a successful attack (i.e. remote code execution etc.). We do not assume to know a probability distribution for this memory, and therefore will not explicitly reason on this part of the code. Instead, note that in this case:

$$\Pr[i \in \Omega(P) \land o \neq \text{crash}] = \Pr[i \in \Omega(P) \land i.\text{payload}[0] \in \text{Valid}]$$

and will depend on the knowledge that the adversary has of Valid.

3.3. Proofs

Let c be a memory-safety exploitation countermeasure. In order to show its effectiveness, given Definition 2, we can leverage the formal semantics of games. To this end, we can proceed in several ways. We can reason directly in the game describing $[P + c]$ and bound the probability $|\Pr[\mathcal{A}^{[P+c]} = i]| \leq \epsilon(n)$ as in the security definition. Alternatively, following the literature in cryptographic proofs, we can proceed to show a bound on $\epsilon(n)$ using game-based transformations, as illustrated in the previous section with the PRF/PRP switching Lemma. The idea of such transformations is to start with the game describing $[P + c]$ and transform it until reaching $[\![P]\!]$. While doing so, we can keep track of the consequences of our transformations (bad flags) to the probability of the event we want to bound. Thanks to the fundamental lemma of game playing [15], we know that:

$$|\Pr[E]_{[P+c]} - \Pr[E]_{[\![P]\!]}| \leq \Pr[F]_{[\![P]\!]}$$

We then leverage on the fact that by definition $\Pr[\mathcal{A}^{[\![P]\!]} = i] = 0$ for $i \in \Omega(P)$ and $[P + c](i) \neq \text{crash}$, and are left to bound the probability of an attacker in distinguishing $[P + c]$ from $[\![P]\!]$.

In the following section we will discuss bounds for various countermeasures and illustrate our technique.

4. Preventing attacks

In this section we show how to model popular probabilistic countermeasures, and how to compute the probability of an attack for countermeasures that aim at preventing attack by crashing attack inputs with high probability.

4.1. Stack Canaries

Stack canaries [21] place a random value between static variables and the return address on the stack. For efficiency, stack canaries have typically a fixed size of one slot on the stack (4 bytes on 32-bit systems) and are not recomputed after each function execution due to performance reasons. However, there is nothing prohibiting an implementation from using larger canary sizes and updating their values after each function execution. For simplicity of illustration, we assume that the size of the canary is potentially arbitrary, and that it is recomputed for each new function execution. In this case:

```
Proc. [P + c](i)
If i ∈ Ω(P) then
    k ←$ {0,1}ⁿ
    ca ← i.payload[0]
    ra ← i.payload[1]
    If ca = k and ra ∈ Valid then
        o ← [M(P, ra)]
    else
        o ← crash
else
    o ← [P](i)
return o
```

Let the event:

$$E = i \in \Omega(P) \land o \neq \text{crash}.$$

Now, it is easy to see that:

$$\begin{aligned}
\Pr[E] &= \Pr[i \in \Omega(P) \land i.\text{payload}[0] = k \\
&\quad \land i.\text{payload}[1] \in \text{Valid}] \\
&\leq \Pr[i \in \Omega(P) \land i.\text{payload}[0] = k] \\
&\leq \Pr[i.\text{payload}[0] = k] = \frac{1}{2^n}
\end{aligned}$$

This follows because the assignment $k \xleftarrow{\$} \{0, 1\}^n$ guarantees that k is an independent and uniformly random variable in the space of bitstrings of length n.

Alternatively, we can use a sequence of game transformations as depicted in Fig 1. We start with the game $[P+c]$ which is semantically equivalent (same output distribution) to $\mathbf{G_1}$ that contains a *bad* flag in case the adversary guesses the canary and a correct return address. $\mathbf{G_1}$ is in turn equivalent up to *bad* to $\mathbf{G_2}$ which will crash even when the adversary guesses the canary and the return address. After removing the bad flag ($\mathbf{G_3}$), merging the redundant branches ($\mathbf{G_4}$) and eliminating deadcode, which are all semantics preserving transformations, we obtain $[\![P]\!]$. Applying the fundamental lemma of game playing we have:

$$
\begin{aligned}
|\Pr[E]_{\mathbf{G_1}} - \Pr[E]_{\mathbf{G_2}}| &\le \Pr[bad = true]_{\mathbf{G_2}} \\
&\le \Pr[i.\mathsf{payload}[0] = k]_{\mathbf{G_2}} \\
&= \frac{1}{2^n}
\end{aligned}
$$

Since in the ideal game $\Pr[E]_{\mathbf{G_2}} = \Pr[E]_{[\![P]\!]} = 0$ we have that $\Pr[E]_{[P+c]} = \Pr[E]_{\mathbf{G_1}} \le \frac{1}{2^n}$. This bound is negligible in n and thus we derive the security proof.

4.1.1. Multiple sampling and single randomization.
Note that, by assumption, the adversary is a PPT. This means that he can at most query the $[P+c]$ oracle a polynomial number of times q, and thus the probability of finding a suitable $i \in \Omega(P)$ such that the execution is not automatically stopped is bounded by $\frac{q}{2^n}$, which is still a negligible function in n. Now, if we modify the game to better model actual implementations, and pass k as a parameter to the execution of P:

```
Proc. G_1
  k ←$ {0,1}^n
```

```
Proc. [P + c](i)
If i ∈ Ω(P) then
    ca ← i.payload[0]
    ra ← i.payload[1]
    If ca = k and ra ∈ Valid then
        o ← [M(P, ra)]
    else
        o ← crash
else
    o ← [P](i)
return o
```

Then the probability of finding $i \in \Omega(P) \land o \ne \mathsf{crash}$ after q queries is given by:

$$
\sum_{j=1}^{q} \frac{1}{2^n - j} \le \frac{q}{2^n - q}
$$

which is still a negligible function in n. Note that in practice, if the value of n is relatively small, then the expected number of queries needed to find a successful guess of the canary is also small, rendering the countermeasure ineffective. This has been pointed out in [14] for ASLR, which as we will see has an identical bound. However, changing the key in every execution of a function has a huge impact in preventing attacks that rely on partial information leakage and *side-channels*, as we will discuss in the following sections.

4.1.2. Concrete bounds.
As discussed before, in practice canaries will have often a 4 byte size and will not be recomputed on each execution (but only on compilation). This gives the bound[1]:

$$
\sum_{j=1}^{q} \frac{1}{2^{32} - j} \le \frac{q}{2^{32} - q}
$$

This certainly is susceptible to brute-force attacks by weak adversaries. To increase security, intuitively, one may increase the size n of the canary. However, such an increase would have a negative impact on performance. If we assume the system size is $s = 32$, then $\lceil \frac{n}{s} \rceil$ stack slots would need to be pop-ed and compared to the canary value on each function return. Moreover, if re-randomization is employed, a larger number of bits must be randomly generated, which also implies an increase in computation effort dependent on the type and implementation of the random number generator.

4.2. ASLR

ASLR [14] is a countermeasure that prevents stack exploitation by randomizing the position of the stack, heap and code in process memory. In the following we assume that the set of valid memory addresses (position of stack, heap and code) has constant size $|\mathsf{Valid}|$, and that the size of memory is variable and of magnitude 2^n. Moreover Π_k is a random permutation of addresses within memory. For simplicity we assume that addresses can be permuted arbitrarily: in practice there will be some constraints on the size of the program, space between heap and stack etc. Also in practice this randomization is done once per system start, however for simplicity we assume this is done after processing any input:

```
Proc. [P + c](i)
k ←$ {0,1}^n
If i ∈ Ω(P) then
    ra ← i.payload[0]
    If ra ∈ Π_k(Valid) then
        o ← [M(P, ra)]
    else
        o ← crash
else
    o ← [P](i)
return o
```

Now, it is easy to see that:

$$
\begin{aligned}
\Pr[i \in \Omega(P) \land o \ne \mathsf{crash}] &\le \Pr[i \in \Omega(P) \\
&\quad \land\ i.\mathsf{payload}[0] \in \Pi_k(\mathsf{Valid})] \\
&\le \Pr[i.\mathsf{payload}[0] \in \Pi_k(\mathsf{Valid})] \\
&= \frac{|\mathsf{Valid}|}{2^n}
\end{aligned}
$$

---

1. Mathematically the denominator of the concrete bounds could thus be negative for certain values of q. However, in practice an attacker will stop before that, because before reaching negative values the fraction will hit 1, which indicates that the attack is successful.
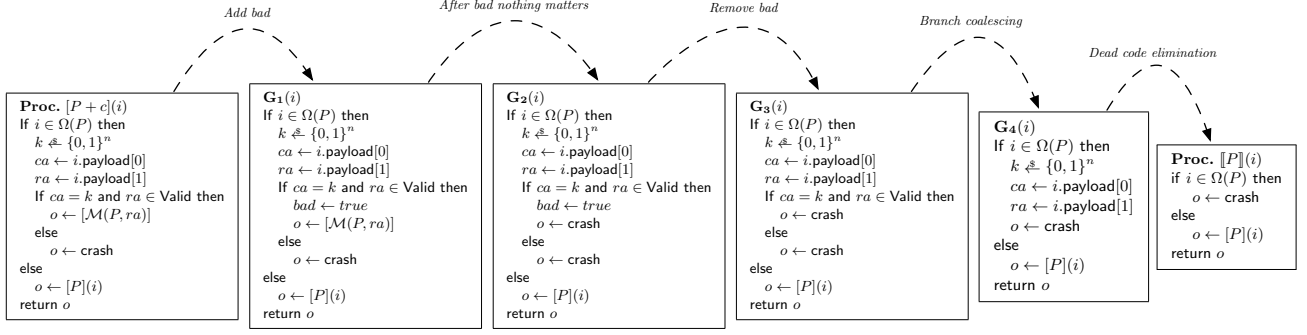
Figure 1. Game transformation for a canary based countermeasure.

This follows because the assignment $k \xleftarrow{\$} \{0,1\}^n$ guarantees that k is an independent an uniformly random variable in the space of bitstrings of length n.

4.2.1. Multiple sampling and single randomization.
Similarly as in the case of canaries, q queries yields a bound of $\frac{q \cdot |\mathsf{Valid}|}{2^n}$, which is still a negligible function in n. Now, if we modify the game to better model actual implementations, and pass k as a parameter to the execution of P:

Proc. G_1
$k \xleftarrow{\$} \{0,1\}^n$

Proc. $[P+c](i)$
If $i \in \Omega(P)$ then
  $ra \leftarrow i.\mathsf{payload}[0]$
  If $ra \in \Pi_k(\mathsf{Valid})$ then
    $o \leftarrow [\mathcal{M}(P, ra)]$
  else
    $o \leftarrow \mathsf{crash}$
else
  $o \leftarrow [P](i)$
return o

Then the probability of finding $i \in \Omega(P) \wedge o \neq \mathsf{crash}$ after q queries is given by:

$$\sum_{j=1}^{q} \frac{|\mathsf{Valid}|}{2^n - j} \leq \frac{q \cdot |\mathsf{Valid}|}{2^n - q}$$

which is still a negligible function in n.

4.2.2. Concrete bounds.
In practice, for 32-bit architectures without virtualization:

$$\sum_{j=1}^{q} \frac{|\mathsf{Valid}|}{2^{32} - j} \leq \frac{q \cdot |\mathsf{Valid}|}{2^{32} - q}$$

To increase security one has to increase the size of the permutation space n. This can be achieved concretely by virtualising memory.

4.3. PointGuard

PointGuard [12] is a countermeasure for BOF attacks which aims to overwrite pointers in any memory location of a process, i.e. stack, heap and static data. PointGuard augments the C compiler by encrypting all pointer values that are stored in memory with a key k and also adds code necessary to decrypt pointer values right before they are loaded into CPU registers (denoted $Decrypt_k$). The only pointer values which are unencrypted during program execution, are those located in CPU registers. Therefore, pointers are never dereferenced directly from process memory, but only via registers. Since registers are not addressable by over-writing pointer values via BOF, the unencrypted pointer values are kept confidential from attackers. The key k is randomly generated on every new process execution and it is never shared outside of the process's address space. If an attacker overwrites the return address on the stack with another address (e.g. of shellcode or of gadgets), then this address will first be decrypted using k and then dereferenced. This will probably lead to crash of the program due to accessing an invalid memory location. To bypass PointGuard an attacker thus has to guess the value of k.

Similarly, to ASLR (see Section 4.2), we assume that the set of valid memory addresses has constant size $|\mathsf{Valid}|$, and that the size of memory is variable and of magnitude 2^n. Also for simplicity we first show the case where k is sampled for every input value passed to the program.

Proc. $[P+c](i)$
$k \xleftarrow{\$} \{0,1\}^n$
If $i \in \Omega(P)$ then
  $ra \leftarrow i.\mathsf{payload}[0]$
  If $Decrypt_k(ra) \in \mathsf{Valid}$ then
    $ra_d \leftarrow Decrypt_k(ra)$
    $o \leftarrow [\mathcal{M}(P, ra_d)]$
  else
    $o \leftarrow \mathsf{crash}$
else
  $o \leftarrow [P](i)$
return o

$Decrypt_k$ is the decryption function (e.g. XOR with key k) used by the PointGuard mechanism. It follows that:

$$\Pr[i \in \Omega(P) \wedge o \neq \mathsf{crash}]$$
$$= \Pr[i \in \Omega(P)$$
$$\wedge Decrypt_k(i.\mathsf{payload}[0]) \in \mathsf{Valid}]$$
$$\leq \Pr[Decrypt_k(i.\mathsf{payload}[0]) \in \mathsf{Valid}]$$
$$= \frac{|\mathsf{Valid}|}{2^n}$$

We can achieve this bound by assuming that an encryption under a random key behaves as a random permutation (and therefore also its inverse).

4.3.1. Multiple sampling and single randomization. Similarly as before, if we modify the game to better model actual implementations, and pass k as a parameter to the execution of P:

<div>

Proc. G_1
$k \xleftarrow{\$} \{0,1\}^n$

</div>

<div>

Proc. $[P + c](i)$
If $i \in \Omega(P)$ then
$\quad ra \leftarrow i.\mathsf{payload}[0]$
$\quad$ If $Decrypt_k(ra) \in \mathsf{Valid}$ then
$\quad\quad ra_d \leftarrow Decrypt_k(ra)$
$\quad\quad o \leftarrow [\mathcal{M}(P, ra_d)]$
$\quad$ else
$\quad\quad o \leftarrow \mathsf{crash}$
else
$\quad o \leftarrow [P](i)$
return o

</div>

Then the probability of finding $i \in \Omega(P) \wedge o \neq \mathsf{crash}$ after q queries is given by:

$$\sum_{j=1}^{q} \frac{|\mathsf{Valid}|}{2^n - j} \leq \frac{q \cdot |\mathsf{Valid}|}{2^n - q}$$

which is still a negligible function in n.

4.3.2. Concrete bounds. For 32-bit architectures such as x86, the size of k is 32 bits. Therefore, the probability of injecting a valid address is identical to that of ASLR (see Section 4.2:

$$\sum_{j=1}^{q} \frac{|\mathsf{Valid}|}{2^{32} - j} \leq \frac{q \cdot |\mathsf{Valid}|}{2^{32} - q}$$

Increasing security in this case is not only related to the size of the memory, but also that of the key k. This means security is not increased by increasing the size of they key beyond the size of a memory pointer. Hence, the size of memory pointers must be increased, e.g. by applying virtualization or switching of a 64-bit architecture.

4.4. ISR

Instruction Set Randomization (ISR) [11] is a countermeasure for code injection attacks via BOFs. Code injection attacks are different from ROP and JOP attacks, because they require the attacker to send executable shellcode as part of their input, while the ROP and JOP attacks require

sending a sequence of gadget addresses and other non-executable data values. Therefore, the code injection attacker can be considered weaker than the ROP or JOP attacker. Nevertheless, relatively recent work [22] has shown that ISR can be combined with stack canaries in order to defend against both attackers. This kind of composite defense follows the *defense in depth* security principle, where a ROP attacker first disables the execution prevention mechanism on code on the stack (e.g. DEP on MS Windows, $W \oplus X$ on Linux), and then executes any code that is injected on the stack. Employing ISR raises the bar for such attacks, therefore, we believe it is interesting to model this protection mechanism.

Similarly to PointGuard, ISR uses encryption via a key k to prevent attacks. Differently from PointGuard it encrypts all the instructions in the code segment, instead of the pointers in all memory segments. Whenever the CPU fetches an instruction to execute, it must be first decrypted using k. Since the key k must be stored in a write-only CPU register which cannot be accessed by attackers, the authors have augmented a CPU emulator with such a register. If an attacker injects shellcode onto the stack via a BOF exploit, this shellcode will be first decrypted (denoted $Decrypt_k$, e.g. XOR with key k) and then executed by the CPU. This will highly likely lead to a crash of the program due to an undefined instruction, i.e. the CPU will not be able to execute the decrypted instruction because it is not part of its *Instruction Set Architecture* (ISA). Therefore, an attacker can bypass ISR by guessing the correct value of k and encrypting the exploit payload with it before injecting it into the vulnerable program. However, ISR can also be bypassed by ROP and JOP attacks [2], [23]. These attacks do not inject code on the stack, instead, they inject pointers to existing code. Since existing code is already encrypted, it will be properly executed by a program which employs ISR as a countermeasure.

<div>

Proc. $[P + c](i)$
$k \xleftarrow{\$} \{0,1\}^m$
If $i \in \Omega(P)$ then
$\quad ra \leftarrow i.\mathsf{payload}[0]$
$\quad$ If $ra \in \mathsf{Valid}$ then
$\quad\quad$ If $Decrypt_k(\mathcal{M}(P, ra)) \in \mathsf{ISA}$ then
$\quad\quad\quad o \leftarrow [\mathcal{M}(P, ra)]$
$\quad\quad$ else
$\quad\quad\quad o \leftarrow \mathsf{crash}$
$\quad$ else
$\quad\quad o \leftarrow \mathsf{crash}$
else
$\quad o \leftarrow [P](i)$
return o

</div>

For the sake of simplicity we assume an ISA with all

instructions having a fixed width of m-bits. It follows that:

$$\begin{aligned}
\Pr[i &\in \Omega(P) \wedge o \neq \mathsf{crash}] \\
&= \Pr[i \in \Omega(P) \\
&\quad \wedge Decrypt_k(\mathcal{M}(P, i.\mathsf{payload}[0]) \in \mathsf{ISA}] \\
&\leq \Pr[Decrypt_k(\mathcal{M}(P, i.\mathsf{payload}[0])) \in \mathsf{ISA}] \\
&= \frac{|\mathsf{ISA}|}{2^m}
\end{aligned}$$

4.4.1. Multiple sampling and single randomization. Similarly as before, if we modify the game to better model actual implementations, and pass k as a parameter to the execution of P:

Proc. G_1
$k \xleftarrow{\$} \{0,1\}^m$

Proc. $[P + c](i)$
If $i \in \Omega(P)$ then
   $ra \leftarrow i.\mathsf{payload}[0]$
   If $ra \in \mathsf{Valid}$ then
      If $Decrypt_k(\mathcal{M}(P, ra)) \in \mathsf{ISA}$ then
         $o \leftarrow [\mathcal{M}(P, ra)]$
      else
         $o \leftarrow \mathsf{crash}$
   else
      $o \leftarrow \mathsf{crash}$
else
   $o \leftarrow [P](i)$
return o

Then the probability of finding $i \in \Omega(P) \wedge o \neq \mathsf{crash}$ after q queries is given by:

$$\sum_{j=1}^{q} \frac{|\mathsf{ISA}|}{2^m - j} \leq \frac{q \cdot |\mathsf{ISA}|}{2^m - q}$$

which is still a negligible function in m.

4.4.2. Concrete bounds. For the 32-bit fixed width ISA we assumed earlier the probability of injecting a valid instruction is:

$$\sum_{j=1}^{q} \frac{|\mathsf{ISA}|}{2^{32} - j} \leq \frac{q \cdot |\mathsf{ISA}|}{2^{32} - q}$$

Increasing security in this case is directly related to the width of the instructions in the ISA and the size of the key k. However, similarly to PointGuard, increasing the size of the key beyond the width of the instructions in the ISA does not increase the security. This can be shown in practice, where the width of instructions between and within an ISA may vary. For instance, in the x86 architecture the smallest instructions are 8-bits, while the longest instructions are 120-bits. This is precisely the insight that Sovarel et al. [24] used to guess the 32-bit key, one-byte at a time, i.e. XOR-ing an 8-bit instruction with a 32-bit key, only uses 8-bits of the key. The search space for guessing a valid 8-bit instruction in

the x86 ISA is reduced from 2^{32} to 2^8. After guessing these 8 bits of the key, the attacker can use a 16-bit instruction to guess the next 8 bits, a.s.o. until the entire 32-bits of the key are guessed. Note that this attack only works in the case of multiple sampling and single randomization, i.e. if the key is changed on every input, then the attack has a much lower success rate. The attack by Sovarel et al. [24] is a prime example of what can go wrong with randomization based defenses if the key size is picked independently of the range of the data items that it must protect.

5. Composition

In practice, in order to raise the bar against attackers, various defense mechanisms are usually stacked together. In the following we discuss examples of composed systems, how to modeal and reason about them using games and their resulting bounds.

5.1. ASLR $\otimes$ Canaries

ASLR and Canaries are commonly used simultaneously in modern systems. We model the composed countermeasure $c = \mathsf{ASLR} \otimes \mathsf{Canaries}$ as follows:

Proc. $[P + c](i)$
If $i \in \Omega(P)$ then
   $k_1 \xleftarrow{\$} \{0,1\}^n$
   $k_2 \xleftarrow{\$} \{0,1\}^m$
   $ca \leftarrow i.\mathsf{payload}[0]$
   $ra \leftarrow i.\mathsf{payload}[1]$
   If $ca = k_1$ then
      If $ra \in \Pi_{k_2}(\mathsf{Valid})$ then
         $o \leftarrow [\mathcal{M}(P, ra)]$
      else
         $o \leftarrow \mathsf{crash}$
   else
      $o \leftarrow \mathsf{crash}$
else
   $o \leftarrow [P](i)$
return o

This modeling is motivated by the fact that when composed, an attacker must bypass both the canary and the ASLR protection in order to avoid a crash.

Now, we can calculate the following bound:

$$\begin{aligned}
\Pr[i \in \Omega(P) \wedge o \neq \mathsf{crash}] &\leq \Pr[i.\mathsf{payload}[0] = k_1 \\
&\quad \wedge i.\mathsf{payload}[1] \in \Pi_{k_2}(\mathsf{Valid})] \\
&\leq \frac{1}{2^n} \cdot \frac{|\mathsf{Valid}|}{2^m}
\end{aligned}$$

This bound indicates that in some cases by combining probabilistic defense mechanisms, we can increase the bounds of attack resistance. However, not all possible combinations will provide an increase. For instance, combining ASLR and PointGuard will not increase the bounds given by PointGuard as we will discuss later in more detail.

5.1.1. Multiple sampling and single randomization. Similarly as in previous examples, if we modify the game to better model actual implementations, and pass k_1 and k_2 as parameters to the execution of P:

Proc. G_1
$k_1 \overset{\$}{\leftarrow} \{0,1\}^n$
$k_2 \overset{\$}{\leftarrow} \{0,1\}^m$

Proc. $[P+c](i)$
If $i \in \Omega(P)$ then
  $ca \leftarrow i.\mathsf{payload}[0]$
  $ra \leftarrow i.\mathsf{payload}[1]$
  If $ca = k_1$ then
    If $ra \in \Pi_{k_2}(\mathsf{Valid})$ then
      $o \leftarrow [\mathcal{M}(P,ra)]$
    else
      $o \leftarrow \mathsf{crash}$
  else
    $o \leftarrow \mathsf{crash}$
else
  $o \leftarrow [P](i)$
return o

Then the probability of finding $i \in \Omega(P) \wedge o \neq \mathsf{crash}$ after q queries trying to guess k_1 and r queries trying to guess k_2 is given by:

$$\left(\sum_{j=1}^{q} \frac{1}{2^n - j}\right) \cdot \left(\sum_{l=1}^{r} \frac{|\mathsf{Valid}|}{2^m - j}\right) \leq \frac{q}{2^n - q} \cdot \frac{r \cdot |\mathsf{Valid}|}{2^m - r})$$

which is still a negligible function in n, m.

5.1.2. Concrete bounds. In practice, for 32-bit architectures without virtualization:

$$\left(\sum_{j=1}^{q} \frac{1}{2^{32} - j}\right) \cdot \left(\sum_{l=1}^{r} \frac{|\mathsf{Valid}|}{2^{32} - l}\right) \leq \frac{q}{2^{32} - q} \cdot \frac{r \cdot |\mathsf{Valid}|}{2^{32} - r}$$

5.2. PointGuard ⊗ ISR

We model the composition $c = \mathsf{PointGuard} \otimes \mathsf{ISR}$ as follows:

Proc. $[P+c](i)$
If $i \in \Omega(P)$ then
  $k_1 \overset{\$}{\leftarrow} \{0,1\}^n$
  $k_2 \overset{\$}{\leftarrow} \{0,1\}^m$
  $ra \leftarrow i.\mathsf{payload}[0]$
  If $Decrypt_{k_1}(ra) \in \mathsf{Valid}$ then
    $ra_d \leftarrow Decrypt_{k_1}(ra)$
    If $Decrypt_{k_2}(\mathcal{M}(P,ra_d)) \in \mathsf{ISA}$ then
      $o \leftarrow [\mathcal{M}(P,ra_d)]$
    else
      $o \leftarrow \mathsf{crash}$
  else
    $o \leftarrow \mathsf{crash}$
else
  $o \leftarrow [P](i)$
return o

This composition has the advantage that it is attack resistant against both code injection attacks, ROP and JOP attacks, because ISR defends against executing injected code and PointGuard defends against overwriting the return

address. We can also perform a similarly attack resistant composition between ISR and ASLR, or ISR and Canaries, or even composing more than two countermeasures. However, we do not show these compositions here due to space restrictions. We can calculate the following bound for Point-Guard ⊗ ISR:

$$\begin{aligned}
&\Pr[i \in \Omega(P) \wedge o \neq \mathsf{crash}] \\
&\quad \leq \Pr[Decrypt_{k_1}(i.\mathsf{payload}[0]) \in \mathsf{Valid} \\
&\quad\quad \wedge \; Decrypt_{k_2}(\mathcal{M}(P, Decrypt_{k_1}(i.\mathsf{payload}[0]))) \in \mathsf{ISA}] \\
&\quad \leq \frac{|\mathsf{Valid}|}{2^n} \cdot \frac{|\mathsf{ISA}|}{2^m}
\end{aligned}$$

Note that the attacker has to not only enter a validly encrypted return address with key k_1 but also to inject a validly encrypted sequence of instructions at that address.

5.2.1. Multiple sampling and single randomization. If we consider implementations of PointGuard and ISR where new keys are generated for each program input, then the model becomes:

Proc. G_1
$k_1 \overset{\$}{\leftarrow} \{0,1\}^n$
$k_2 \overset{\$}{\leftarrow} \{0,1\}^m$

Proc. $[P+c](i)$
If $i \in \Omega(P)$ then
  $ra \leftarrow i.\mathsf{payload}[0]$
  If $Decrypt_{k_1}(ra) \in \mathsf{Valid}$ then
    $ra_d \leftarrow Decrypt_{k_1}(ra)$
    If $Decrypt_{k_2}(\mathcal{M}(P,ra_d)) \in \mathsf{ISA}$ then
      $o \leftarrow [\mathcal{M}(P,ra_d)]$
    else
      $o \leftarrow \mathsf{crash}$
  else
    $o \leftarrow \mathsf{crash}$
else
  $o \leftarrow [P](i)$
return o

In this case, the probability of finding an input that violates memory safety $i \in \Omega(P)$, after q queries trying to guess k_1 and r queries trying to guess k_2, is equal to:

$$\left(\sum_{j=1}^{q} \frac{|\mathsf{Valid}|}{2^n - j}\right) \cdot \left(\sum_{l=1}^{r} \frac{|\mathsf{ISA}|}{2^m - j}\right) \leq \frac{q \cdot |\mathsf{Valid}|}{2^n - q} \cdot \frac{r \cdot |\mathsf{ISA}|}{2^m - r},$$

which is a negligible function in n, m.

5.2.2. Concrete bounds. For 32-bit architectures where the ISA has a fixed width, the bounds are equal to:

$$\left(\sum_{j=1}^{q} \frac{|\mathsf{Valid}|}{2^{32} - j}\right) \cdot \left(\sum_{l=1}^{r} \frac{|\mathsf{ISA}|}{2^{32} - j}\right) \leq \frac{q \cdot |\mathsf{Valid}|}{2^{32} - q} \cdot \frac{r \cdot |\mathsf{ISA}|}{2^{32} - r},$$

Composition	n-bit Architecture	32-bit Architecture	64-bit Architecture	128-bit Architecture
ASLR⊗PointGuard	$\frac{q\cdot\lvert\mathsf{Valid}\rvert}{2^n-q}$	1	2^{-23}	2^{-87}
ASLR⊗ISR	$\frac{q\cdot\lvert\mathsf{Valid}\rvert}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{ISA}\rvert}{2^n-r}$	2^{-10}	2^{-75}	2^{-203}
PointGuard⊗ISR	$\frac{q\cdot\lvert\mathsf{Valid}\rvert}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{ISA}\rvert}{2^n-r}$	2^{-10}	2^{-75}	2^{-203}
Canary⊗ASLR	$\frac{q}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{Valid}\rvert}{2^n-r}$	2^{-22}	2^{-87}	2^{-215}
Canary⊗PointGuard	$\frac{q}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{Valid}\rvert}{2^n-r}$	2^{-22}	2^{-87}	2^{-215}
Canary⊗ISR	$\frac{q}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{ISA}\rvert}{2^n-r}$	2^{-26}	2^{-91}	2^{-219}
ASLR⊗PointGuard⊗ISR	$\frac{q\cdot\lvert\mathsf{Valid}\rvert}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{ISA}\rvert}{2^n-r}$	2^{-10}	2^{-75}	2^{-203}
Canary⊗ASLR⊗PointGuard	$\frac{q}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{Valid}\rvert}{2^n-r}$	2^{-22}	2^{-87}	2^{-215}
Canary⊗ASLR⊗ISR	$\frac{q}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{Valid}\rvert}{2^n-r}\cdot\frac{t\cdot\lvert\mathsf{ISA}\rvert}{2^n-t}$	2^{-42}	2^{-139}	2^{-331}
Canary⊗PointGuard⊗ISR	$\frac{q}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{Valid}\rvert}{2^n-r}\cdot\frac{t\cdot\lvert\mathsf{ISA}\rvert}{2^n-t}$	2^{-42}	2^{-139}	2^{-331}
Canary⊗ASLR⊗PointGuard⊗ISR	$\frac{q}{2^n-q}\cdot\frac{r\cdot\lvert\mathsf{Valid}\rvert}{2^n-r}\cdot\frac{t\cdot\lvert\mathsf{ISA}\rvert}{2^n-t}$	2^{-42}	2^{-139}	2^{-331}

TABLE 1. CONCRETE BOUNDS FOR COUNTERMEASURE COMPOSITIONS

5.3. Overview of compositions

According to AT&T[2] the current overall average network latency is 33 ms and no inter-USA state latency is faster than 4 ms. Therefore, in the worst case we assume that a remote attacker to be able to submit attack requests (queries) with a speed of 1 query/ms. Given this optimistic attack speed of 1 query/ms we also assume that after a total number of 2^{25} queries, a network administrator or an intrusion prevention system will block the attacker. This translates into over 9 days of continuous queries sent by the attacker, which again aims to depict a worst case scenario.

We do not explicitly model the remaining possible compositions of countermeasures presented in Section 4. However, in Table 1 we provide an overview of both the general formulas and the concrete bounds of the remaining compositions, for 32-bit, 64-bit and virtualized 128-bit architectures, where q, r and t represent the number of guesses for the different keys of the composed countermeasures. Note that in order to obtain the values in this table we also considered the fact that the size of the ISA is 2^{12} in the worst case, because even in the Intel x86 ISA we have 3683 instructions, if we consider all possible mnemonics and operand types [25]. Finally, we assume the size of $\lvert\mathsf{Valid}\rvert$ is 2^{16}, because according to Follner et al. [26], the number of ROP gadgets in a programs range from a few hundred in programs as small as bzip2, to a little under 60.000 ROP gadgets for programs such as GCC. For many practical purposes the number of ROP gadgets can be a useful estimation of $\lvert\mathsf{Valid}\rvert$, however, if one wants to use a sound upper bound the total number of instructions can be counted. Since our paper assumes that ISAs have a fixed width equal to the size of one word, the size of $\lvert\mathsf{Valid}\rvert = 2^{16}$ corresponds roughly to 2^{18}, 2^{19} and 2^{20} byte programs (ca. 256 KBs, 512 KBs and 1 MBs) for 32-, 64- and 128-bit architectures, respectively.

Since these countermeasures target different parts of the program (e.g. instructions for ISR, pointers for PointGuard and stack memory for Canaries), or are applied at different points of the program lifetime (e.g. compile-time for ISR, PointGuard and Canaries or load-time for ASLR), the order in which these countermeasures is fixed, i.e. the order of applying the countermeasures cannot be changed.

Note that any composition involving PointGuard and ASLR takes the same bounds of either one of the countermeasures. This is due to the fact that both countermeasures force attackers to guess a valid memory address. On the other hand, ISR and Canaries force attackers to guess the key used to encrypt instructions and the canary value, respectively. Therefore, the lowest bound obtainable from any composition involves ISR, Canaries and either PointGuard or ASLR.

6. Side-channel attacks

An attacker can use side-channel attacks to get information about a secret key by exploiting physical observations such as time, power consumption, noise, etc. For instance, if re-randomization is not employed, canaries are susceptible to brute force attacks using crashing as a side-channel [27]. The key insight employed by this attack is that any integer number of bytes can be written passed the bounds of a buffer, (not necessarily a multiple of the machine word size). So far, our modeling assumes that the adversary overwrites both the canary and the return address entirely. However, in practice s/he may choose to only partially overwrite the canary, to learn it byte by byte, which is computationally much faster than guessing all bytes at once. The attacker uses the fact that the program crashes, as an oracle for his guesses.

In other words, this means that the attacker can overwrite the vulnerable buffer, up to and including only the first byte of the canary. If this overwritten byte is incorrect, then the program crashes, otherwise, the first byte of the canary was guessed correctly and the attacker can move on to guessing the second byte of the canary. This side-channel attack reduces the search space from 2^k to $2^8\times\lceil\frac{n}{8}\rceil$, which is a linear function of the key size n. To be able to capture this

attack, we model canaries more accurately in the following way:

$$
\boxed{\begin{array}{l} \textbf{Proc. } G_1 \\ \quad k \overset{\$}{\leftarrow} \{0,1\}^n \end{array}}
$$

$$
\boxed{\begin{array}{l}
\textbf{Proc. } [P+c](i) \\
\text{If } i \in \Omega(P) \text{ then} \\
\quad ca \leftarrow i.\mathsf{payload}[0] \\
\quad ra \leftarrow i.\mathsf{payload}[1] \\
\quad \text{If } ow(k, ca) = k \text{ and } ra \in \mathsf{Valid} \text{ then} \\
\quad\quad o \leftarrow [\mathcal{M}(P, ra)] \\
\quad \text{else} \\
\quad\quad o \leftarrow \mathsf{crash} \\
\text{else} \\
\quad o \leftarrow [P](i) \\
\text{return } o
\end{array}}
$$

where $ow(a,b)$ returns a string of length $|ow(a,b)| = |a|$ and overwrites the first $|b|$ bytes of a with the corresponding values in b, leaving the remaining values identical to the corresponding positions in a:

$$
ow(a,b)[i] = \begin{cases} b[i] & \text{if } 0 < i \leq |b| \\ a[i] & \text{otherwise} \end{cases}
$$

Therefore, if the the first entry of the payload ($i.\mathsf{payload}[0]$) is 1 byte in size, then $ow(k, ca)$ will only overwrite the first byte of the canary k.

6.1. Reasoning about side-channel leakage

We can formally plug the leakage λ learned by an adversary to our bounds by decreasing the entropy on the key (originally n bits) by λ bits. For instance, if there is a side-channel leaking λ bits of information on the ASLR key to the adversary after q_1 observations, and then an attacker performs q_2 subsequent observations, then our bound becomes:

$$
\Pr[i \in \Omega(P) \wedge o \neq \mathsf{crash}] \leq \frac{q_2 \cdot |\mathsf{Valid}|}{2^{n-\lambda} - q_2}
$$

Depending on how λ increases with respect to q and n, this could have practical relevance for concrete scenarios.

6.1.1. Concrete bounds. Take the side-channel attack on stack canaries presented at the beginning of Section 6. If we consider $q_1 = 256 = 2^8$, i.e. 256 guesses made by the attacker on the first byte of the canary, then we can consider $\lambda = 8$-bits, since the attacker has guessed the first byte of the canary, reducing the search space by that amount. Note that in practice he needs in average just 128 tries to guess the first byte, here we consider the worst case. Since the size of the canary can be assumed to be 4 bytes (i.e. $n = 32$) as we discussed in Section 4.1.2, the bound on a successful attack becomes:

$$
\Pr[i \in \Omega(P) \wedge o \neq \mathsf{crash}] \leq \frac{1}{2^{24}}
$$

which is at least 4 orders of magnitude higher than $\frac{1}{2^{32}}$. If afterwards the attacker performs q_2 queries, this bounds becomes:

$$
\Pr[i \in \Omega(P) \wedge o \neq \mathsf{crash}] \leq \frac{q_2}{2^{24} - q_2}
$$

In general, it is difficult to foresee and prevent all possible side-channels, in particular if they consider time or power consumption. However we can prevent many side-channels that are reflected directly in the values sent to the attacker (such as crashes and termination, or memory leaks such as Heartbleed) by relying on simultaneous executions of the victim program, as presented in Section 6.2.

6.2. Probabilistic countermeasures and replicas

The use of replicas (similarly to Secure Multi-Execution, SME [19]) can prevent such information leakages. For instance, in the case of canaries, if there are at least two replicas with a different canary each, the probability of an attacker overriding both with the same values is the probability of both having the same bytes: If the canary consists of 1 byte, then the probability of that byte (pseudo-randomly generated) being the same in both replicas r_0 and r_1 is $\Pr[cb^{r_0} = cb^{r_1}] = \frac{1}{2^8}$. If the canary consists of n bytes, then $\Pr[\bigwedge_{i=0}^{n-1} cb_i^{r_0} = cb_i^{r_1}] = \frac{1}{2^{8n}}$.

This reflects precisely the idea behind SME, instead of having a direct mapping for each input of the system whether it has high or low-level security, the security level is given by the nature of the input given. Valid outputs are considered of a high security level and the crash output due to inconsistent replicas is considered as having a low security level. Then, for a valid input (high security level), all its valid outputs are calculated and given to the user (since they also have a high security level). However, if its invalid (low security level), with high probability the program will crash (which will be the result of the low-level security execution.

We can generalize this idea as depicted in Fig. 2. For any countermeasure relying on a secret key, we execute two copies of the program for different keys k and k'. Before sending the resulting output to the remote adversary, we compare the resulting outputs $o = [P+c]_k[i]$ and $o' = [P+c]_{k'}[i]$. We assume that both programs use otherwise the same seeds for randomness and if they interact with the system they read the same values (such as the system's clock, values of files, sockets etc.). If they are different, it means that the output is dependent on the only different input for the two copies: the secret keys. In this case, we send a crash message to the adversary and reset the program's copies with two fresh keys. However, *legal*

Figure 2. SME inspired replicas to close value-based side-channels on $[P+c]$.

inputs, that is $j \notin \Omega$ will be in general independent from random keys used by the probabilistic countermeasure, thus this check is transparent for such inputs.

6.3. Hardening probabilistic countermeasures with replicas

We note that, not only it is possible to effectively close many side-channels using SME, but we can also make the probabilistic countermeasures even more effective. In particular, we make exploitation attacks impossible in the following way.

Lemma 1. Let $[P + c]$ be a probabilistic countermeasure against remote attacks. Let a prerequisite for exploitation that the attacker produces a payload that matches a random key. Formally, $i.\mathsf{payload}[0] = k$ for a deterministic oracle $\mathsf{payload}$. Then, for a concrete instance of such a countermeasure using key k, an input $i \in \Omega(P)$ such that $[P + c]_k[i] \neq \mathsf{crash}$ must crash another instance of the countermeasure with a key $k' \neq k$: $[P + c]_{k'}[i] = \mathsf{crash}$.

The proof follows simply by the fact that $k \neq k'$ and that the payload oracle is deterministic. From this lemma, it follows:

Corollary 1. Let $\mathsf{SME}([P + c])$ be an implementation of $[P + c]$ as described in Section 6.2 and in Fig. 2, that is, it executes two instances of $[P + c]$ in parallel with random keys k, k' such that $k \neq k'$, and outputs crash whenever $[P + c]_k \neq [P + c]_{k'}$. Then:

$$\Pr[\mathcal{A}^{\mathsf{SME}([P+c])} = i] = 0$$

where $i \in \Omega(P)$ and $\mathsf{SME}([P + c])(i) \neq \mathsf{crash}$.

The proof of the corollary follows immediately from the lemma above. Surprisingly, the keys k and k' can be fully leaked to the attacker in this setting, since (mathematically) it is impossible to produce an input that would bypass both instances of the countermeasure.

Although this is a strong result abstractly speaking, from the point of view of practical implementations it suffers from several drawbacks. First, it is questionable whether executing a program twice will be faster than embedding a deterministic defense such as a run-time monitor. On the other hand, depending on how the parallel execution is implemented, it could be that an attacker could craft an attack on one instance that disables the check on the second instance. In principle it is possible to avoid such attacks for instance using virtualization.

One can argue however that this idea is interesting in settings where for legacy reasons it is not possible to recompile the code with memory safety checks and where there are no hard time constraints on the execution, but a thorough efficiency and efficacy analysis is left to future work.

7. Related work

The question on the effectiveness of probabilistic countermeasures against memory-safety vulnerabilities has been a matter of discussion in the literature for the past decade. In [14], Schacham et al. discuss the effectiveness of Address Space-Layout Randomization. They noted that in 32 bits architectures, the entropy achieved by ASLR only slightly slows down attackers, while causing performance overheads. They also discuss the increase in security by re-randomization, but different from us, they do this informally and only for ASLR.

Pucella et al. [28] propose to treat the effectiveness of program diversification (i.e. using probabilistic transformations) as a probabilistic dynamic type-checking problem. They illustrate their approach on a C-like language and ASLR. Their approach is formal but involves the modeling of low-level details of programs, such as pointers, whereas our approach abstracts away from concrete programs and focus on countermeasures. Also, they do not explicitly give bounds on the probability of an attack for a given program transformation.

Abadi et al. [29] cast the problem of reasoning on the effectiveness of ASLR as full-abstraction problem. Different from Pucella et al., they consider concrete probability bounds on attackers, and abstract away from malicious inputs by considering arbitrary execution contexts. To this end, they also construct exemplary high and low level languages, and limit themselves to the analysis of ASLR.

Berger et al. [20] propose to randomize the heap layout and allocation strategy for increased resilience against memory management bugs that are triggered accidentally. This line of work has been extended [30] by considering attackers that deliberately exploit vulnerabilities (using for instance heapspray attacks). Probabilistic security bounds are then derived by reasoning on the proposed countermeasures and successful attack events. Such bounds, although mathematically justified, are obtained informally (i.e. without using a formal reasoning language or framework).

In our examples we have emphazised single vs. multiple key-sampling for randomized countermeasures and

discussed its impact on the bounds. Such a fine-grained re-randomization for many countermeasures and memory layouts has been systematically implemented by Curtsinger and Berger [31].

To the best of our knowledge, we are the first to propose game-based proofs for memory-safety countermeasures and to apply the generic approach to several countermeasures and their composition.

8. Conclusions

In this paper we have presented an approach to reason about probabilistic countermeasures against memory-safety vulnerabilities in a rigorous way using concepts from cryptographic proofs. We have shown that our modeling is applicable for a wide range of countermeasures and their composition. Moreover we have shown how to close certain side-channels using replicas, similarly as in Secure Multi-Execution (SME). Surprisingly, using this technique also hardens significantly several probabilistic countermeasures, by making exploitation (theoretically) infeasible. In future work, we plan to model other diversity inspired countermeasures from the literature and reason about their guarantees, such as DieHarder [30]. Moreover we plan to use existing tool support to rigorously develop computer-aided proofs of our bounds.

References

[1] Y. Younan, W. Joosen, and F. Piessens, "Runtime countermeasures for code injection attacks against c and c++ programs," *ACM Computing Surveys (CSUR)*, vol. 44, no. 3, p. 17, 2012.

[2] H. Shacham, "The geometry of innocent flesh on the bone: Return-into-libc without function calls (on the x86)," in *Proceedings of the 14th ACM conference on Computer and communications security*. ACM, 2007, pp. 552–561.

[3] A. Bittau, A. Belay, A. Mashtizadeh, D. Mazieres, and D. Boneh, "Hacking blind," in *Security and Privacy (SP), 2014 IEEE Symposium on*. IEEE, 2014, pp. 227–242.

[4] K. Z. Snow, F. Monrose, L. Davi, A. Dmitrienko, C. Liebchen, and A.-R. Sadeghi, "Just-in-time code reuse: On the effectiveness of fine-grained address space layout randomization," in *Security and Privacy (SP), 2013 IEEE Symposium on*. IEEE, 2013, pp. 574–588.

[5] Z. Durumeric, J. Kasten, D. Adrian, J. A. Halderman, M. Bailey, F. Li, N. Weaver, J. Amann, J. Beekman, M. Payer *et al.*, "The matter of heartbleed," in *Proceedings of the 2014 Conference on Internet Measurement Conference*. ACM, 2014, pp. 475–488.

[6] M. Abadi, M. Budiu, U. Erlingsson, and J. Ligatti, "Control-flow integrity," in *Proceedings of the 12th ACM conference on Computer and communications security*. ACM, 2005, pp. 340–353.

[7] B. Zeng, G. Tan, and Ú. Erlingsson, "Strato: A retargetable framework for low-level inlined-reference monitors." in *Usenix Security*, 2013, pp. 369–382.

[8] E. Göktas, E. Athanasopoulos, H. Bos, and G. Portokalidis, "Out of control: Overcoming control-flow integrity," in *2014 IEEE Symposium on Security and Privacy*. IEEE, 2014, pp. 575–589.

[9] S. Bhatkar, D. C. DuVarney, and R. Sekar, "Address obfuscation: An efficient approach to combat a broad range of memory error exploits." in *USENIX Security*, vol. 3, 2003, pp. 105–120.

[10] C. Cowan, S. Beattie, R. F. Day, C. Pu, P. Wagle, and E. Walthinsen, "Protecting systems from stack smashing attacks with stackguard," in *Linux Expo*. Citeseer, 1999.

[11] G. S. Kc, A. D. Keromytis, and V. Prevelakis, "Countering code-injection attacks with instruction-set randomization," in *Proc. of the 10th ACM conference on Computer and communications security*, 2003.

[12] C. Cowan, S. Beattie, J. Johansen, and P. Wagle, "Pointguard tm: protecting pointers from buffer overflow vulnerabilities," in *Proceedings of the 12th conference on USENIX Security Symposium*, vol. 12, 2003, pp. 91–104.

[13] Y. Weiss and E. G. Barrantes, "Known/chosen key attacks against software instruction set randomization," in *Computer Security Applications Conference, 2006. ACSAC'06. 22nd Annual*. IEEE, 2006, pp. 349–360.

[14] H. Shacham, M. Page, B. Pfaff, E.-J. Goh, N. Modadugu, and D. Boneh, "On the effectiveness of address-space randomization," in *Proc. of the 11th ACM conference on Computer and communications security*. ACM, 2004, pp. 298–307.

[15] M. Bellare and P. Rogaway, "The security of triple encryption and a framework for code-based game-playing proofs," in *Advances in Cryptology-EUROCRYPT 2006*. Springer, 2006, pp. 409–426.

[16] O. Goldreich, "The foundations of modern cryptography," in *Modern Cryptography, Probabilistic Proofs and Pseudorandomness*. Springer, 1999, pp. 1–37.

[17] V. Ganesh, S. Banescu, and M. Ochoa, "Short paper: The meaning of attack-resistant systems," in *Proceedings of the 10th ACM Workshop on Programming Languages and Analysis for Security, PLAS@ECOOP 2015, Prague, Czech Republic, July 4-10, 2015*, 2015, pp. 49–55.

[18] G. Barthe, B. Grégoire, and S. Zanella Béguelin, "Formal certification of code-based cryptographic proofs," *ACM SIGPLAN Notices*, vol. 44, no. 1, pp. 90–101, 2009.

[19] D. Devriese and F. Piessens, "Noninterference through secure multi-execution," in *Security and Privacy (SP), 2010 IEEE Symposium on*. IEEE, 2010, pp. 109–124.

[20] E. D. Berger and B. G. Zorn, "Diehard: Probabilistic memory safety for unsafe languages," in *Proc. of the 2006 ACM SIGPLAN Conference on Programming Language Design and Implementation*. ACM, 2006, pp. 158–168. [Online]. Available: http://doi.acm.org/10.1145/1133981.1134000

[21] P. Wagle and C. Cowan, "Stackguard: Simple stack smash protection for gcc," in *Proceedings of the GCC Developers Summit*. Citeseer, 2003, pp. 243–255.

[22] A. Papadogiannakis, L. Loutsis, V. Papaefstathiou, and S. Ioannidis, "Asist: architectural support for instruction set randomization," in *Proceedings of the 2013 ACM SIGSAC conference on Computer & communications security*. ACM, 2013, pp. 981–992.

[23] T. Bletsch, X. Jiang, V. W. Freeh, and Z. Liang, "Jump-oriented programming: a new class of code-reuse attack," in *Proceedings of the 6th ACM Symposium on Information, Computer and Communications Security*. ACM, 2011, pp. 30–40.

[24] A. N. Sovarel, D. Evans, and N. Paul, "Where's the feeb? the effectiveness of instruction set randomization." in *Usenix Security*, 2005.

[25] S. Heule, "How many x86-64 instructions are there anyway?" https://stefanheule.com/blog/how-many-x86-64-instructions-are-there-anyway/, accessed: 2017-01-12.

[26] A. Follner, A. Bartel, and E. Bodden, "Analyzing the gadgets towards a metric to measure gadget quality," *arXiv preprint arXiv:1605.08159*, 2016.

[27] A. 'pi3' Zabrocki, "Scraps of notes on remote stack overflow exploitation," November 2010.

[29] M. Abadi and G. D. Plotkin, "On protection by layout randomization," *ACM Transactions on Information and System Security (TISSEC)*, vol. 15, no. 2, p. 8, 2012.

[30] G. Novark and E. D. Berger, "Dieharder: securing the heap," in *Proceedings of the 17th ACM conference on Computer and communications security*. ACM, 2010, pp. 573–584.

[31] C. Curtsinger and E. D. Berger, "Stabilizer: statistically sound performance evaluation," in *ACM SIGARCH Computer Architecture News*, vol. 41, no. 1. ACM, 2013, pp. 219–228.

[28] R. Pucella and F. B. Schneider, "Independence from obfuscation: A semantic framework for diversity," *Journal of Computer Security*, vol. 18, no. 5, pp. 701–749, 2010.