

ECE351 Course Outline, Winter 2014

Vijay Ganesh

January 6, 2014

vganesh@uwaterloo.ca

1 Introduction

Welcome to ECE351! This is an introductory compilers course. The calendar description lists a number of topics:

Programming paradigms, compilation, interpretation, virtual machines. Lexical analysis, regular expressions and finite automata. Parsing, context-free grammars and push-down automata. Semantic analysis, scope and name analysis, type checking. Intermediate representations. Control flow. Data types and storage management. Code generation.

We will cover many of the topics listed above. We will do so by closely following the book “Crafting a Compiler” (ISBN 13: 978-0-13-606705-4). Please purchase a copy asap. For the lab part of the course, a manual authored by Professor Derek Rayside will be used. It is available on the course website.

1.1 Course Website

All relevant course materials will be available from the course website linked from <https://ece.uwaterloo.ca/~vganesh/teaching.html>

1.2 Course Objective

The objectives of the course are two-fold: First, to teach students the theoretical aspects of compilers and language design, and second, to enable students to build a practical rudimentary compiler of their own. The course is structured as follows: We will have one lecture every week that focuses on theory, and another lecture on lab.

This course is very hands on, i.e., you will write a circuit synthesizer and a circuit simulator for circuits described in a subset of VHDL. A circuit synthesizer is essentially a compiler for circuits. The lab notes describe these exercises in detail.

Through the lab you will learn two important skills: First, you will understand how language design can be seen as a general form of problem solving. Second, you will learn how to “implement” a language by constructing an appropriate compiler.

Often language design and problem solving are seen as distinct activities. The point of the lab is to show that language design, in many ways, is the most general way to solve a problem. By most general solution I mean that your solution solves an entire class of problems,

VHDL is a very complicated hardware description language to specify circuits. The language for which you will develop a synthesizer is a very small subset of VHDL.

not merely one problem or an instance of a problem. For example, suppose that you need to multiply two matrices. The least general, and most direct, thing you could do is to actually multiply the matrices. More generally, you could write a program that multiplies two arbitrary matrices, and then use that program to multiply your two matrices of interest. Most generally, you could design a language that not only allows you to write programs to multiply matrices but also allows you to express programs to do more complex mathematical operations. Your language could be something similar to a suitable subset of Matlab, then write a program in that language that requests your two matrices be multiplied, then execute that program with the interpreter/compiler you just developed. In order to solve problems at this level of abstraction you need to understand the ideas listed in the course description, and you will need to improve your programming skills.

We will use the term *transformer* or *translator* to describe a program that reads and writes structured text. As you will see, any sufficiently large transformer is actually comprised of many small transformers: the transformation is broken down into steps that typically pass through one or more intermediate forms before the final output is produced. These intermediate forms are typically simpler to process mechanically and less convenient for human use than the original input language.

Our VHDL transformers will make use of a variety of intermediate forms. In the first half of the term we will write transformers just for these intermediate languages. Only the second half of the term will we process VHDL input. The lab notes give a detailed schedule of the weekly labs, as well as a summary of some necessary computer skills.

Solve an entire class of problems by designing a language to describe those problems.

Good engineers do a cost/benefit analysis before designing and building. Good engineers also have the skills to design a range of solutions.

A compiler is a particular kind of transformer that reads source code and produces assembly or machine code.

For example, the GNU Compiler Collection (GCC) has an intermediate form called RTL (Register Transfer Language). C/C++/Java/Fortran/etc. are first translated to RTL, then optimizations are performed on the RTL, and finally assembly code is generated for some particular chip.

2 Instructor/TAs Offices and Contact, Classroom and Lab Location

Instructor:	Vijay Ganesh	DC2530	vganesh@uwaterloo.ca
Lab Instructor:	Mohamed Hassan		m49hassa@uwaterloo.ca
TAs:	Wenzhu Man		wman@uwaterloo.ca
	Leo Passos		lpassos@gsd.uwaterloo.ca
Lecture Date and Time:	MF 10:00AM–11:20AM		RCH103
Tutorial:	Thu 4:30PM–5:20PM	E2-3353	for help with labs
Midterm:	none		
Lab reports:	Thursdays 11:59PM	Git	every week (except midterm week)
Lab:	TWTh 10am & 3pm	E2-3353	http://sun5.vlsi.uwaterloo.ca/~ecepc/TimeTables/
Files:	Git Repo		https://ecegit.uwaterloo.ca/
Discussion:	Piazza		https://piazza.com/uwaterloo.ca/winter2014/ece351

Files will be distributed and collected only through the Git Repo. Each student will have their own private repository. We will have shared repositories for skeleton code and test cases. Your Nexus credentials (username and password) should give you appropriate access to the repositories.

3 Collaboration

Interaction is an essential part of learning for most people. We are going to try a novel structure in ECE351 to facilitate learning collaboratively and honestly. We here define different levels of collaboration that correspond to different maximum grades. Your grade for a lab will be the lower of your earned grade or the cap. For each lab you will declare your level of collaboration. While we give you the flexibility of declaring a different level of collaboration for each lab, we expect that your level will remain largely constant throughout the course. The levels are:

Individual: You discuss the labs with fellow students, perhaps making sketches on a blackboard/whiteboard. No written artifacts leave the discussion and get transferred to your code, i.e., your code is completely written by you.

Partner: You and a partner collaborate. You may look at each other's code. Each student is still expected to do all of the typing on his or her computer. Partners are expected to be at a roughly equivalent skill level.

Mentor: A mentor teaches a protégé. The mentor is expected to teach the protégé at the protégé's computer, and only the protégé operates the computer. The protégé does not look at the mentor's computer.

At every level of collaboration, every student is expected to physically key in every character that he or she commits. Students are not permitted to share electronic files with each other, or with students who have taken this course in the past, except as facilitated by the instructors using the course version control system.

4 Reference Material

The textbook for the course is *Crafting a Compiler*¹. Past offerings of this course have recommended *Modern Compiler Implementation in Java*² or *Programming Language Pragmatics*³. Those are both good books, but we think *Crafting a Compiler* is a better fit for this course. All three books are available at the library.

The intellectual content is widely available on the internet. In the past some students have reported that they have found the internet to be a big place and they have had some difficulty figuring out which parts of it are relevant to the course. Perhaps your searching skills are better.

The tools that we will use are documented largely online.⁴

Pre-lab material on the computing environment is excluded from these collaboration restrictions. By all means, please help each other out getting the computing environment working.

Skeleton code provided by course staff is excluded from these collaboration restrictions. You may look at the skeleton code on a computer collaboratively *before* you or your partner start writing your solutions and still declare *verbal* collaboration.

cap 100%

cap 90%

mentor cap 100%
protégé cap 80%

Course staff are excluded from the collaboration caps, i.e., you may ask questions of the staff, they may look at your code, *etc.*, without that imposing a cap on your mark.

We will share test cases this way, for example.

¹ Charles N. Fischer, Ron K. Cytron, and Richard J. LeBlanc, Jr. *Crafting a Compiler*. Addison-Wesley, Reading, Mass., 2010

² Andrew W. Appel and Jens Palsberg. *Modern Compiler Implementation in Java*. Cambridge University Press, 2004

³ Michael L. Scott. *Programming Language Pragmatics*. Morgan Kaufmann, 3 edition, 2009

⁴ parboiled.org

5 Marking Scheme

Labs	50%
Final Exam	50%
Total	100%

- a. You must pass the final exam in order to pass the course.
- b. You must earn at least 50% of the available lab marks in order to pass the course.
- c. Late lab submissions are ignored. Note that you may commit+push your work as many times as you want before the deadline. We will mark the code that is on the main branch of the repository at the deadline.
- d. We will remark up to three labs, at your request, at the end of the term. The value to be gained from this remarking is capped at 5% of the overall grade.
- e. There might be bonus marks available for exceptional class participation: submitting patches, answering forum questions, mentoring other students, *etc.*

We recommend that you commit+push frequently — at least at the end of every work session.

This provision is to give you some credit for fixing bugs in earlier labs, or allows you to completely miss one lab. The labs are cumulative.

6 Classroom Style

The projector will be used to display the text book, the course notes, slides, the lab manual, and the code.

The board will be used to primarily to work through problems. All problems will be taken from a text book or the course notes, so you will not have to write them down yourself.

Speak up if you do not understand what is being discussed in class. Chances are what you do not understand is some background material that you weren't actually taught in your previous courses.

7 University Policies

Academic Integrity: In order to maintain a culture of academic integrity, members of the University of Waterloo community are expected to promote honesty, trust, fairness, respect and responsibility.

<http://uwaterloo.ca/academicintegrity/>

Grievance: A student who believes that a decision affecting some aspect of his/her university life has been unfair or unreasonable may have grounds for initiating a grievance. When in doubt please be certain to contact the department's administrative assistant who will provide further assistance.

Policy 70, Student Petitions and Grievances, §4, <http://secretariat.uwaterloo.ca/Policies/policy70.htm>

Discipline: A student is expected to know what constitutes academic integrity to avoid committing an academic offence, and to take responsibility for his/her actions. A student who is unsure whether an action constitutes an offence, or who needs help in learning how to avoid offences (e.g., plagiarism, cheating) or about rules for group work/collaboration should seek guidance from the course instructor, academic advisor, or the undergraduate Associate Dean.

<http://uwaterloo.ca/academicintegrity/>

For information on categories of offences and types of penalties, students should refer to Policy 71, Student Discipline, <http://secretariat.uwaterloo.ca/Policies/policy71.htm>

For typical penalties check Guidelines for the Assessment of Penalties, <http://secretariat.uwaterloo.ca/guidelines/penaltyguidelines.htm>

Appeals: A decision made or penalty imposed under Policy 70 (Student Petitions and Grievances) (other than a petition) or Policy 71 (Student Discipline) may be appealed if there is a ground. A student who believes he/she has a ground for an appeal should refer to Policy 72 (Student Appeals).

<http://secretariat.uwaterloo.ca/Policies/policy70.htm>

<http://secretariat.uwaterloo.ca/Policies/policy71.htm>

<http://secretariat.uwaterloo.ca/Policies/policy72.htm>

Note for Students with Disabilities: The Office for Persons with Disabilities (OPD) collaborates with all academic departments to arrange appropriate accommodations for students with disabilities without compromising the academic integrity of the curriculum. If you require academic accommodations to lessen the impact of your disability, please register with the OPD at the beginning of each academic term.

OPD: Needles Hall, Room 1132

Plagiarism Detection Software: The instructors may, at their discretion, use plagiarism detection software.

e.g., TurnItIn.com

References

- [1] Andrew W. Appel and Jens Palsberg. *Modern Compiler Implementation in Java*. Cambridge University Press, 2004.
- [2] Charles N. Fischer, Ron K. Cytron, and Richard J. LeBlanc, Jr. *Crafting a Compiler*. Addison-Wesley, Reading, Mass., 2010.
- [3] Michael L. Scott. *Programming Language Pragmatics*. Morgan Kaufmann, 3 edition, 2009.